

# Uncertainty

Prof Tom Drummond

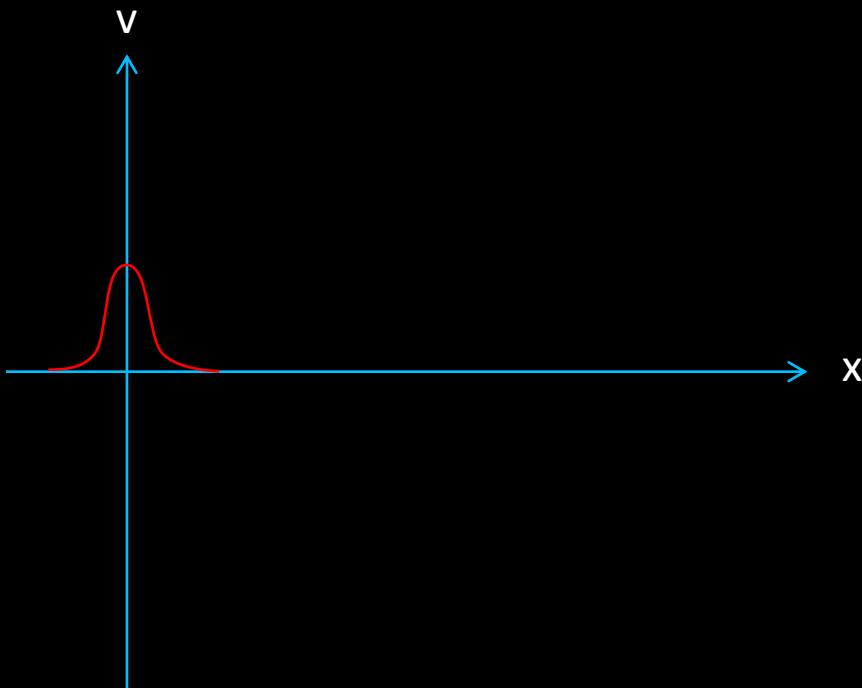
## Claim 1:

It's more important to know what you don't  
know than to get the right answer

# Kalman Filtering

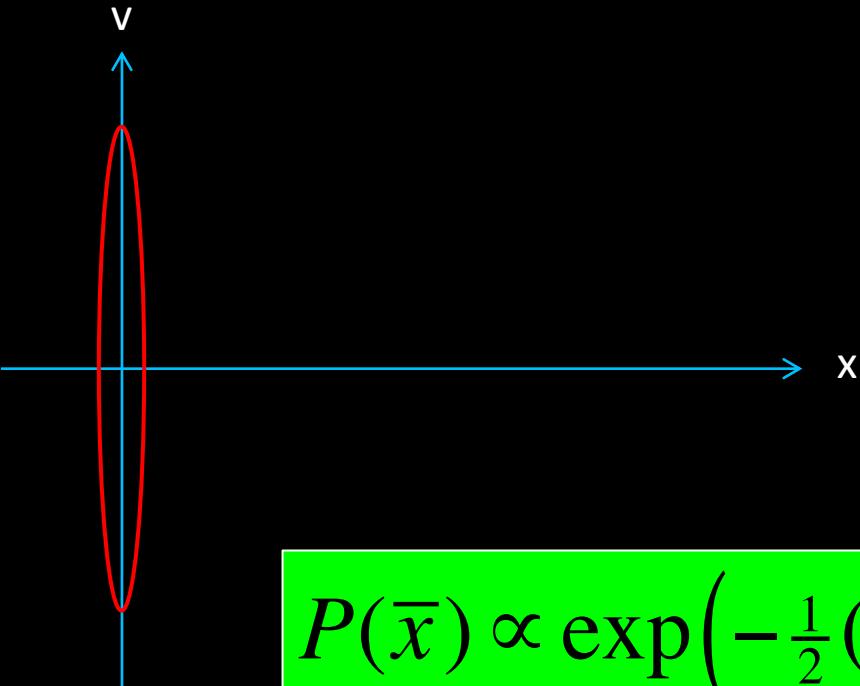
Simple example with 1D position measurement and 1D velocity.

At time  $t=0$ , we make an uncertain measurement of  $x$  and get the value zero.



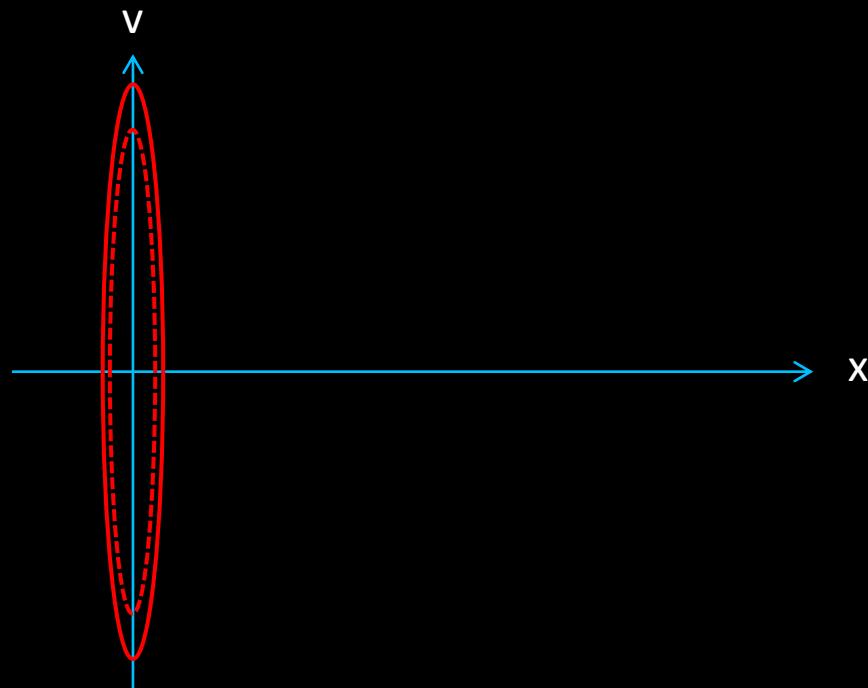
# Kalman Filtering

Initialize filter with position measurement and big uncertainty in velocity:



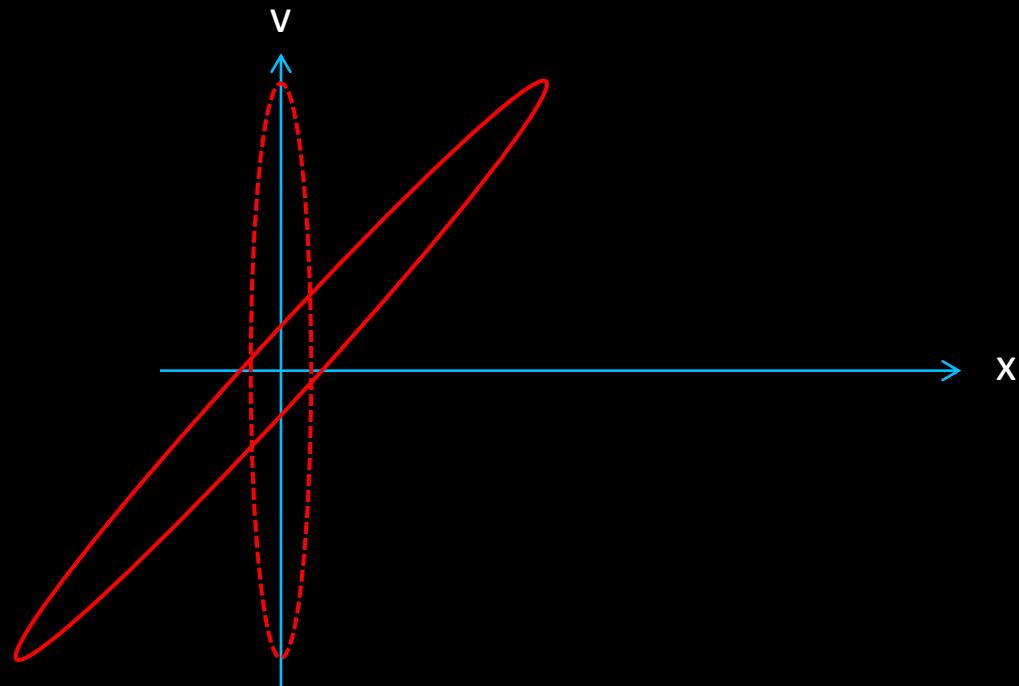
# Kalman Filtering

Add process noise...



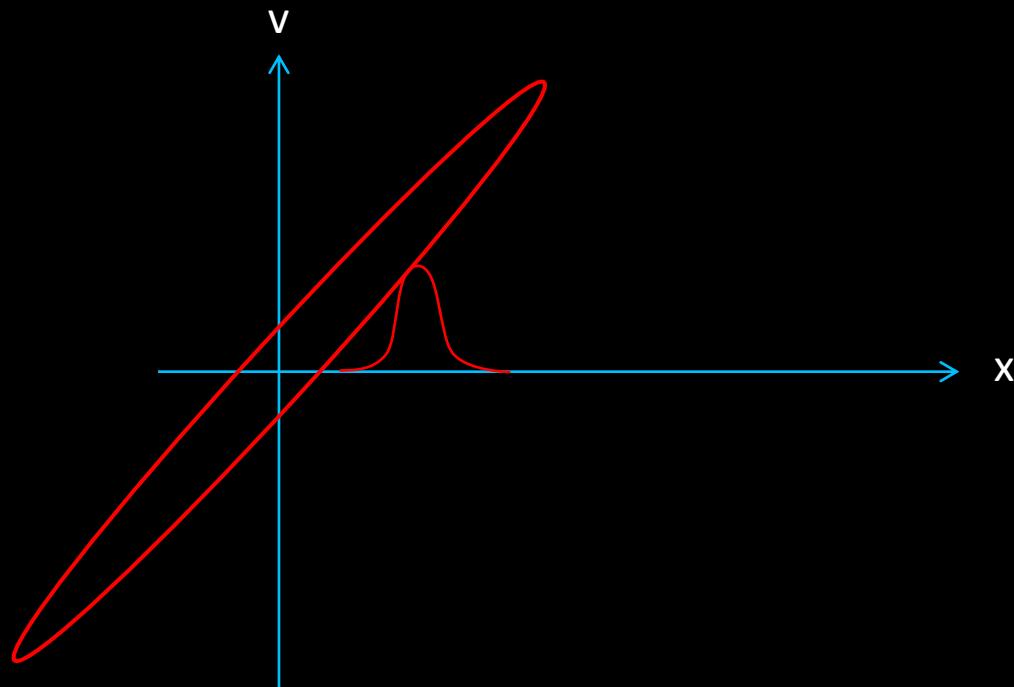
# Kalman Filtering

Time passes – position and velocity uncertainties become correlated:



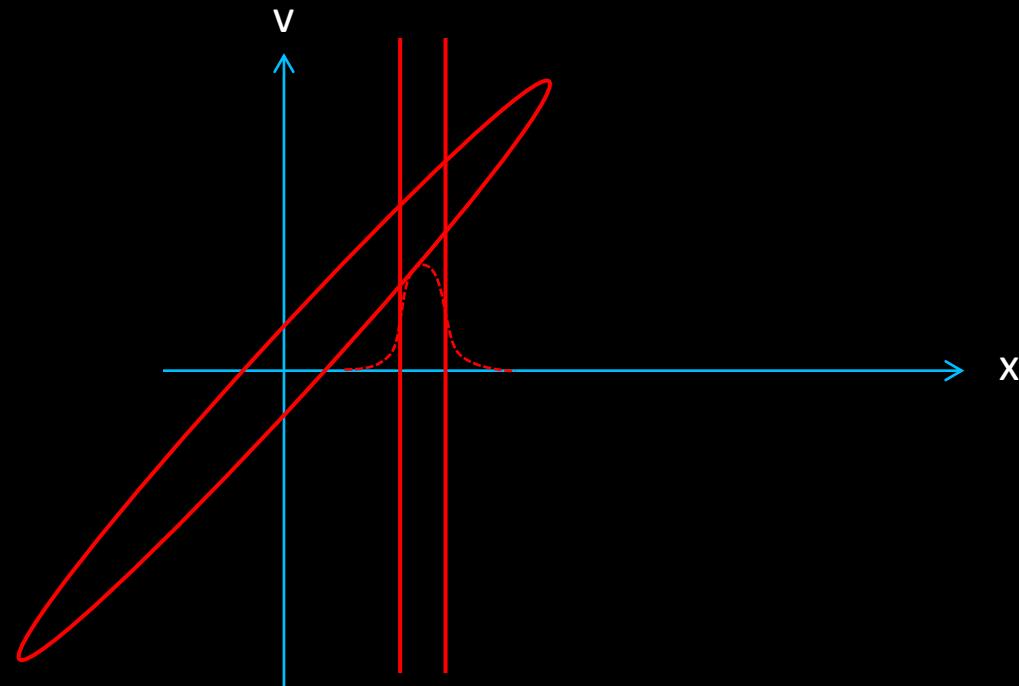
# Kalman Filtering

Take a second measurement of position



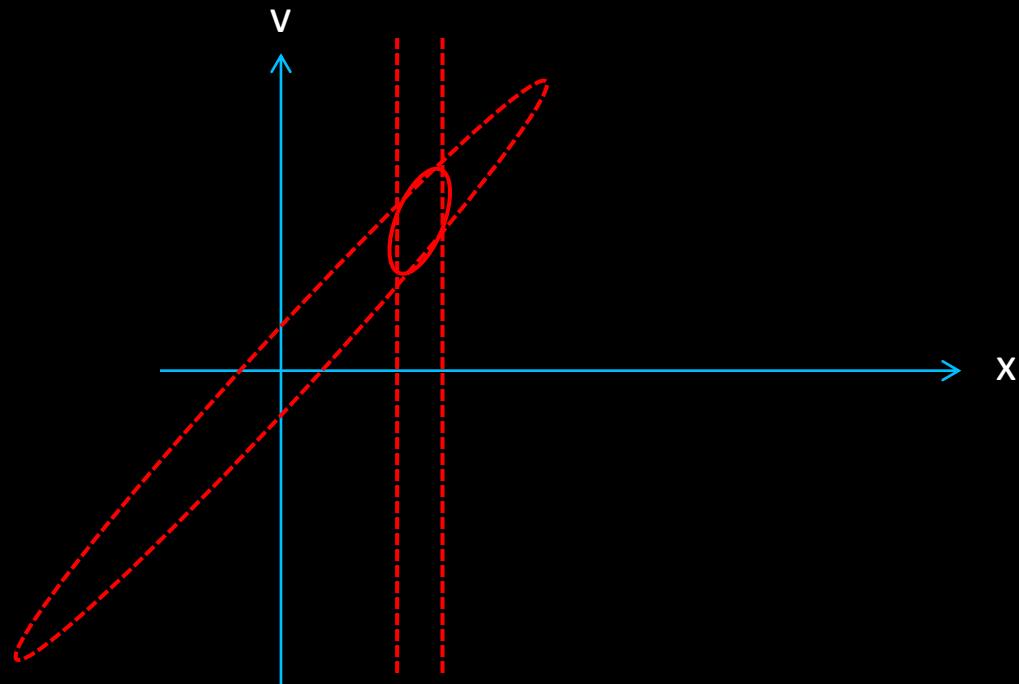
# Kalman Filtering

This corresponds to a measurement with some variance in position and infinite variance in velocity



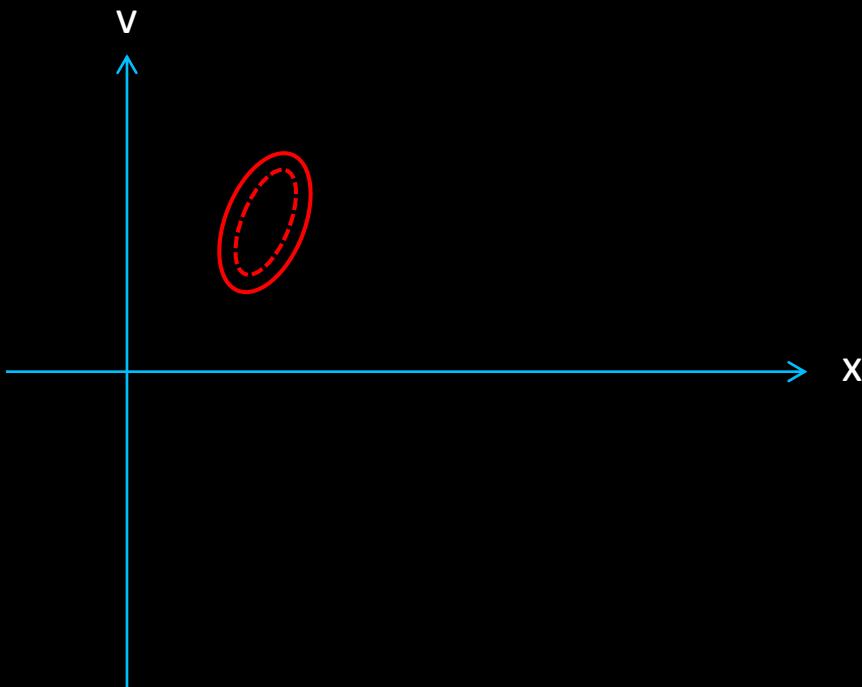
# Kalman Filtering

Combine the prior with the new measurement information to obtain a posterior estimate



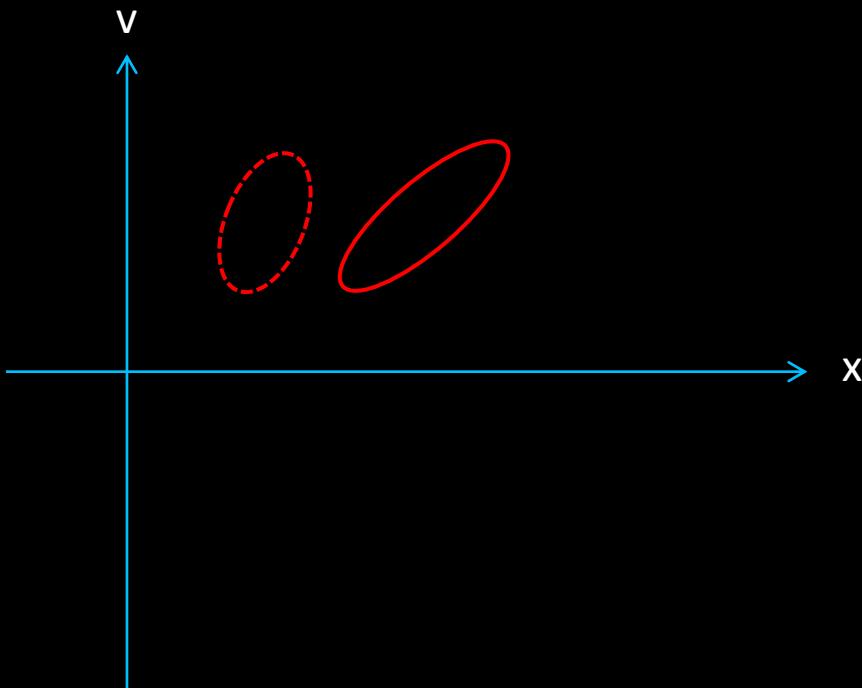
# Kalman Filtering

Add process noise...



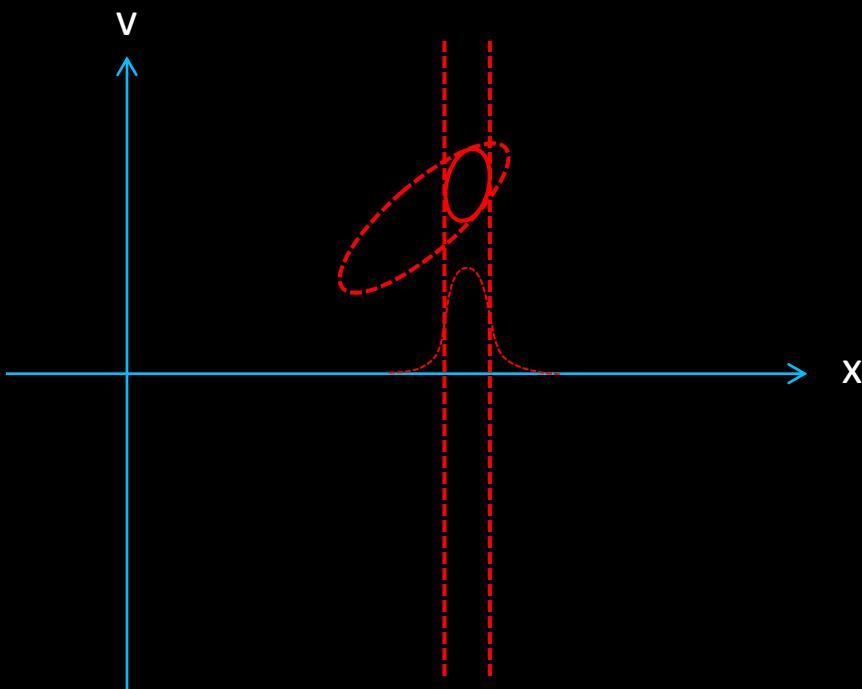
# Kalman Filtering

More time passes...

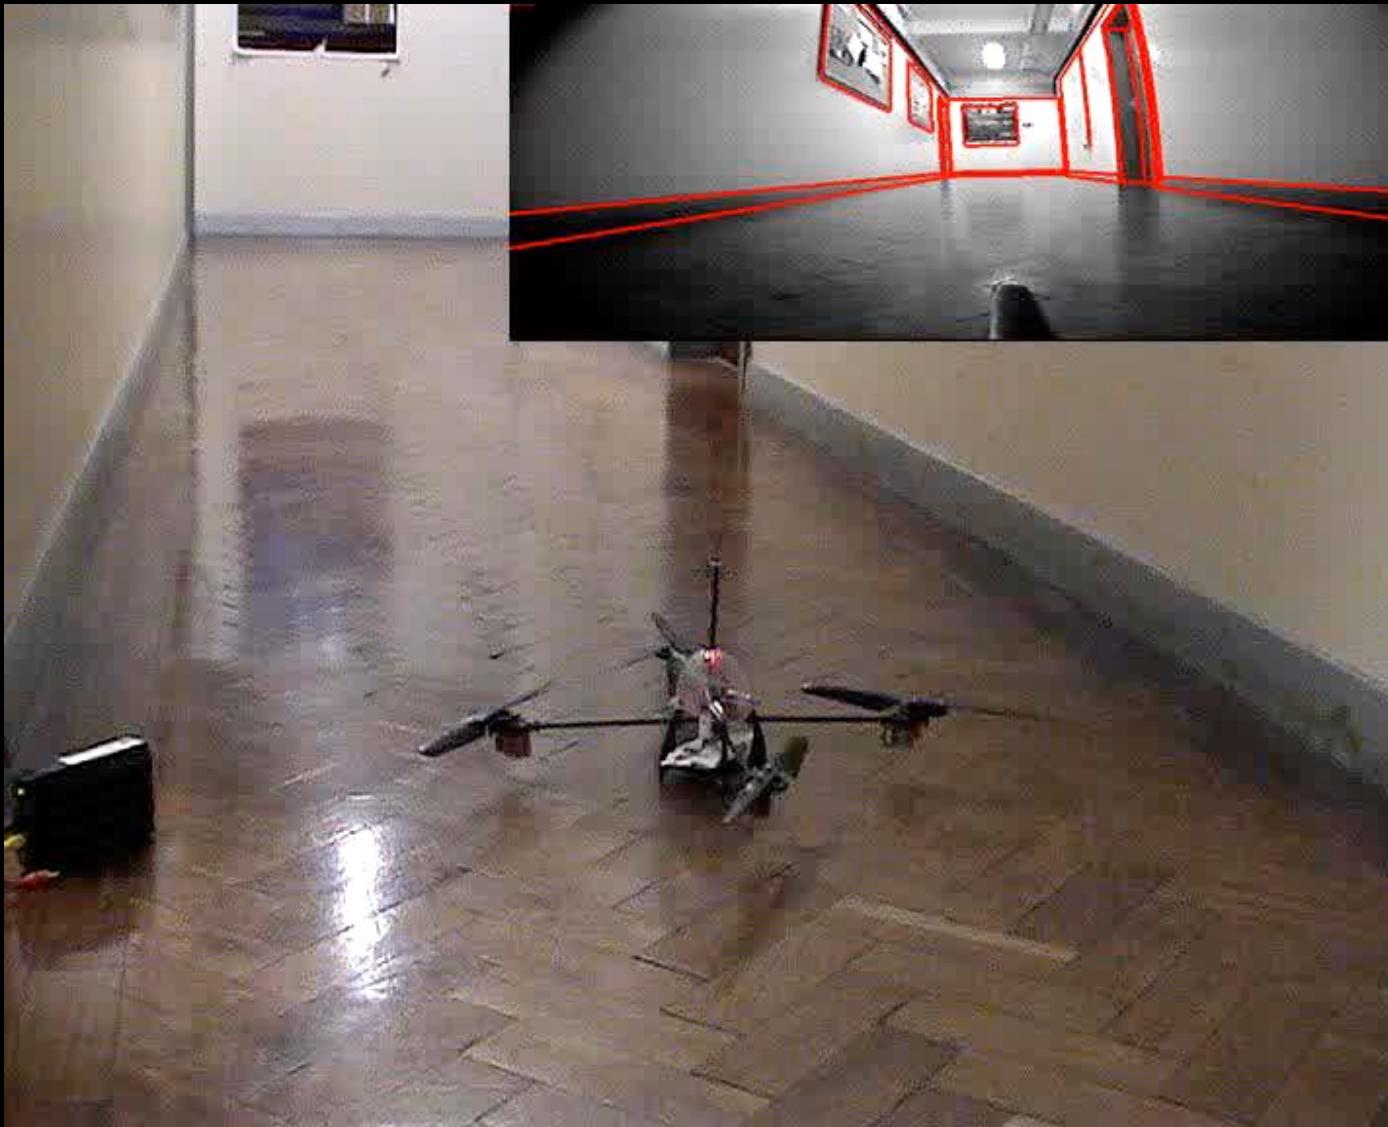


# Kalman Filtering

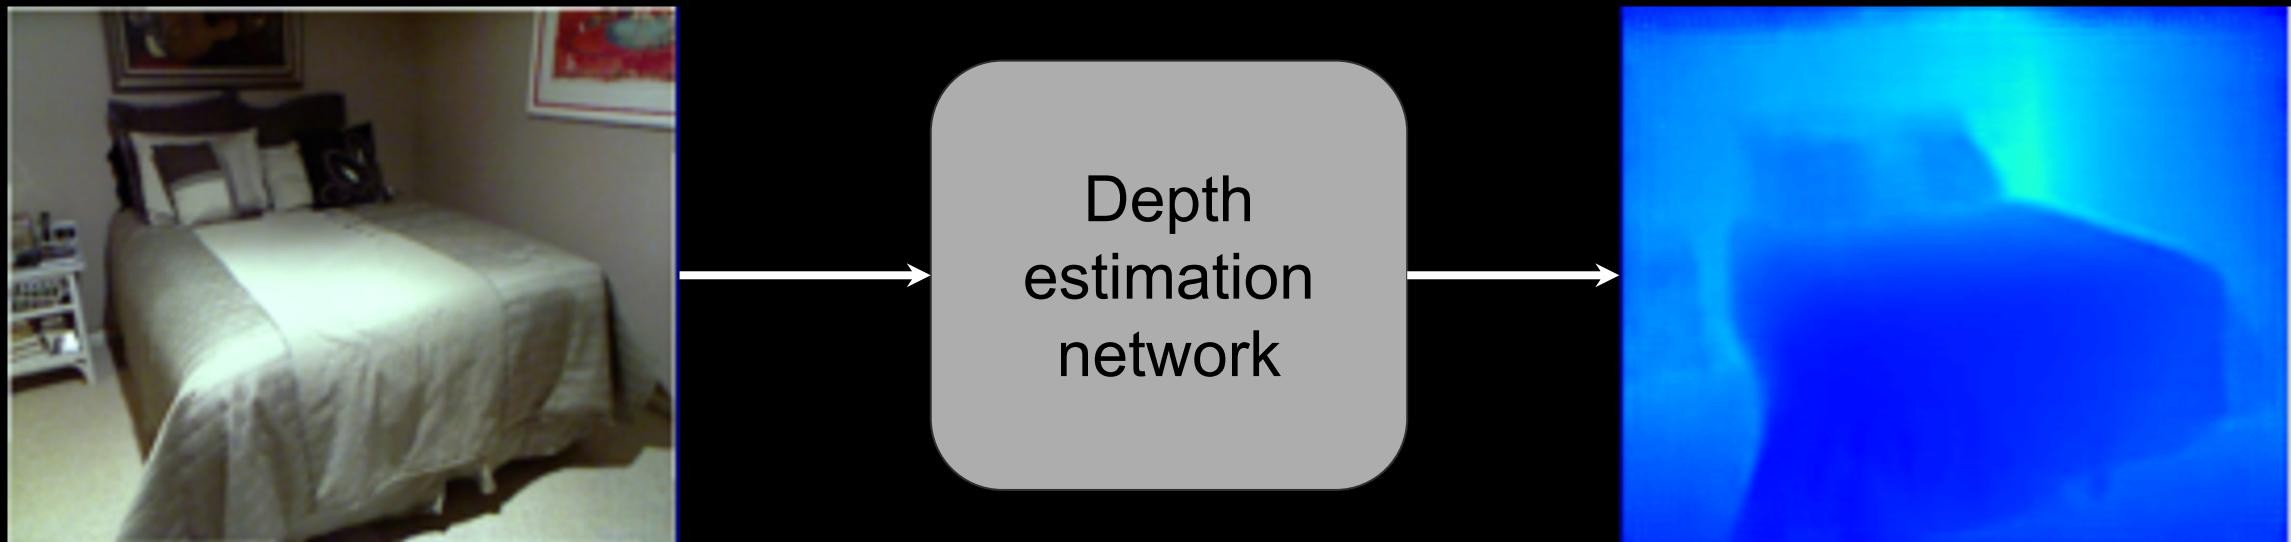
Make a third measurement:



# What if the errors are not Gaussian?

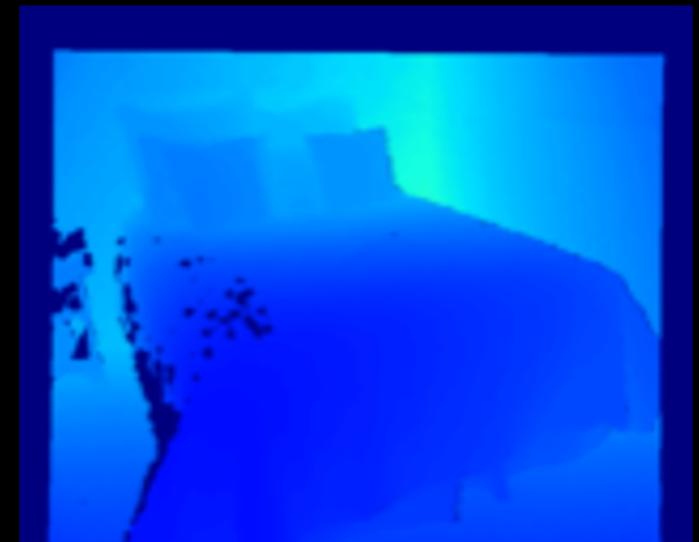


# Can use learning to estimate depth from a single image



Minimise  $L_2$  loss to ground truth.

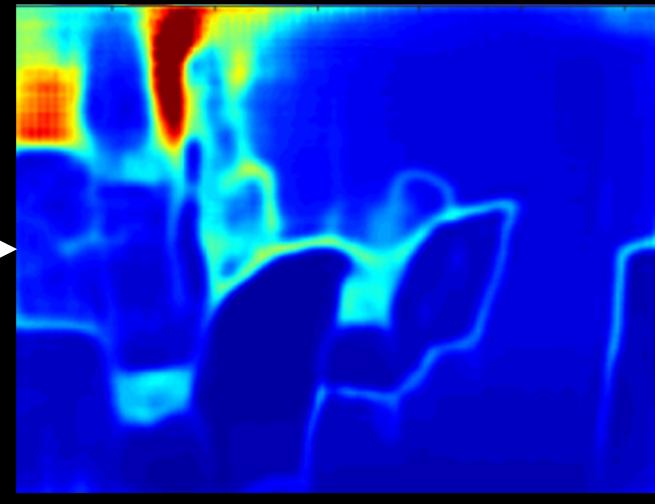
Q: Should we estimate  $z$ ,  $1/z$  or  $\log(z)$ ?



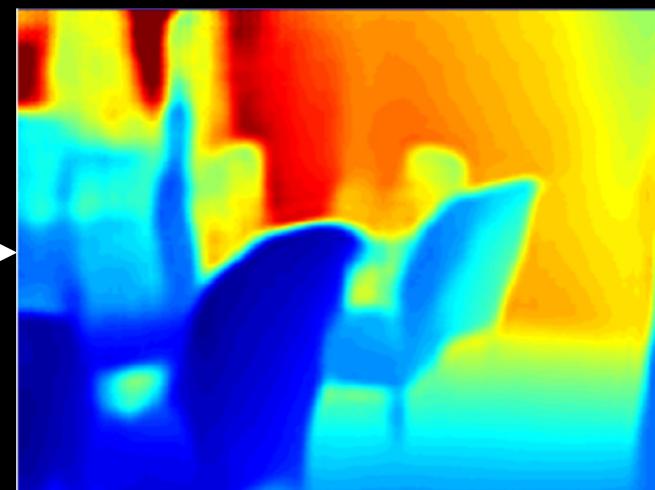
# Better to learn uncertainty as well as depth estimates



Depth  
estimation  
network



Uncertainty ( $\sigma$ )

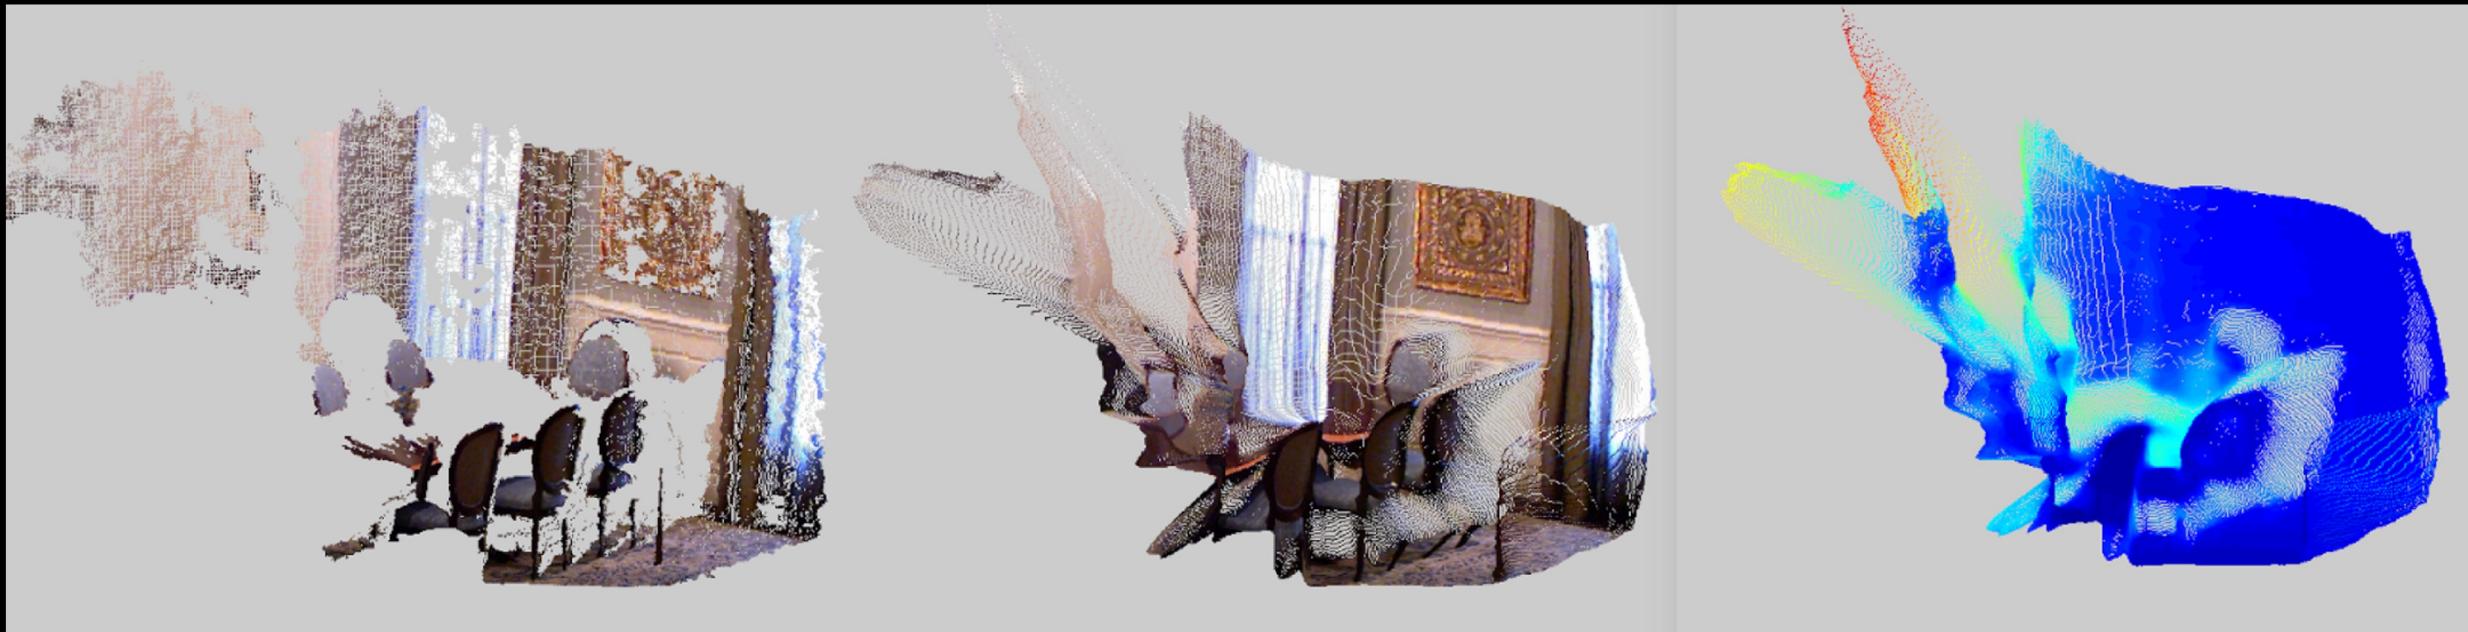


Depth ( $\mu$ )

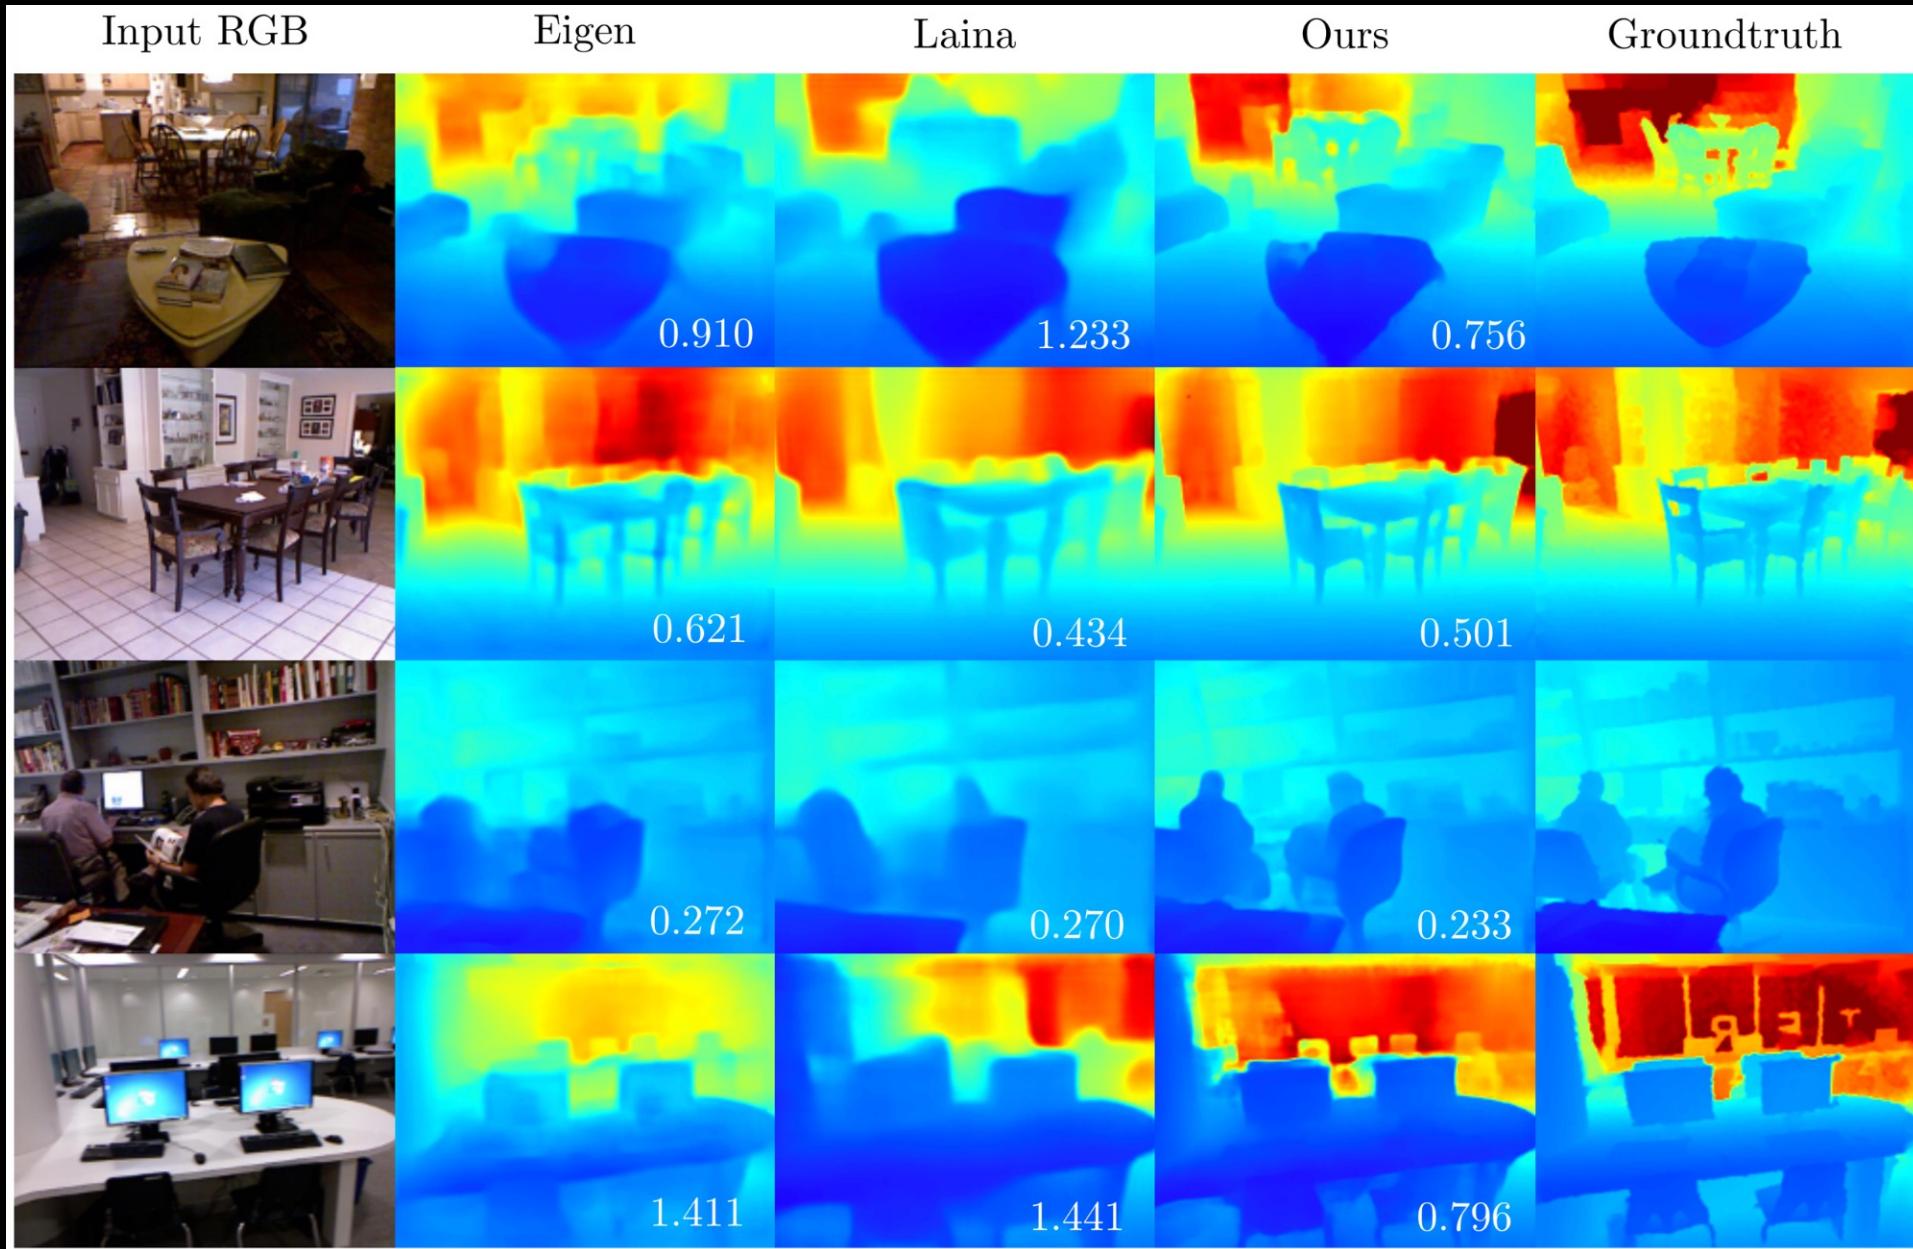
Q: How can we learn  $\sigma$  if we only have ground truth for  $\mu$ ?

# Estimating uncertainty in learning

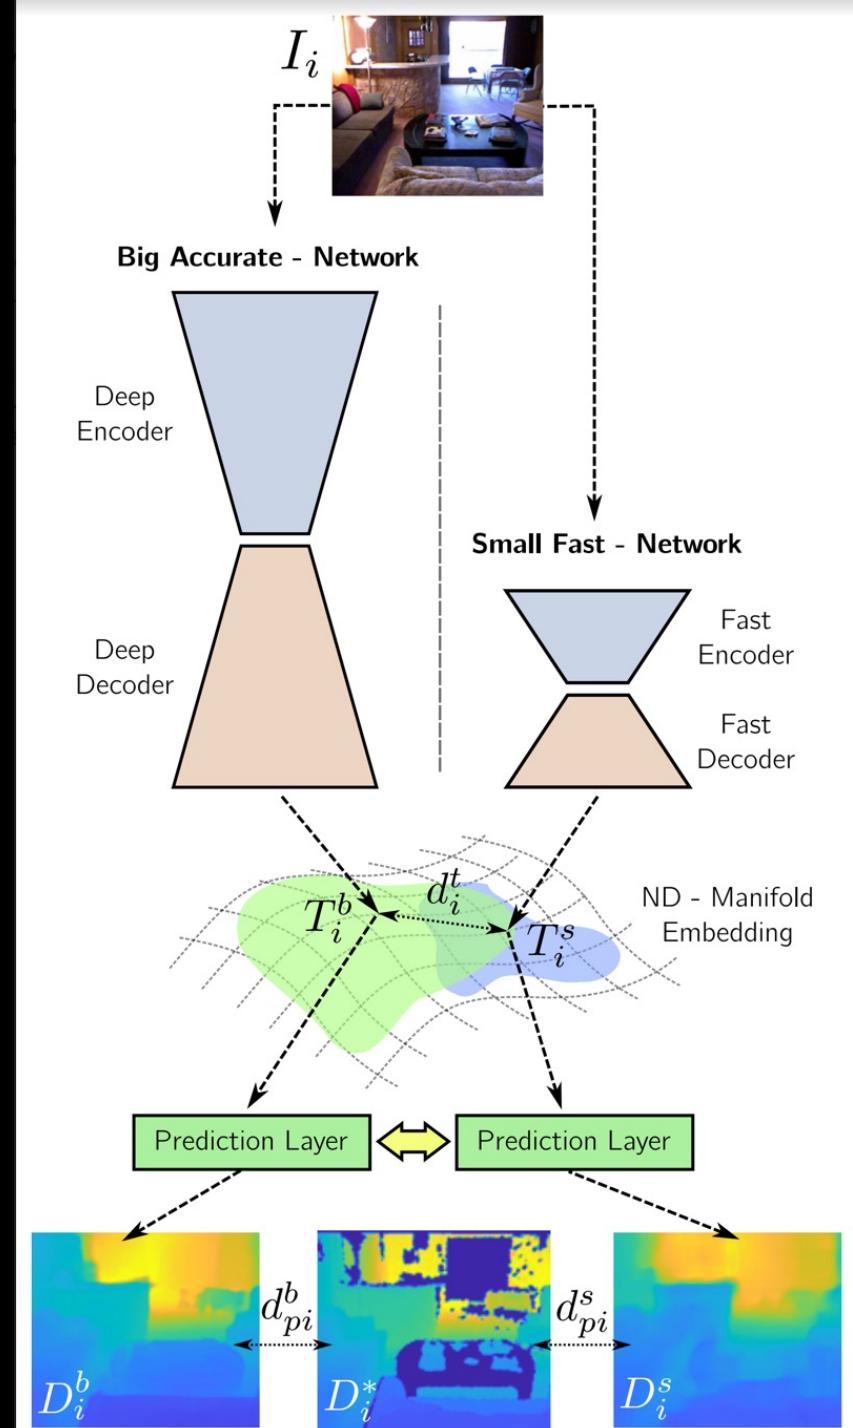
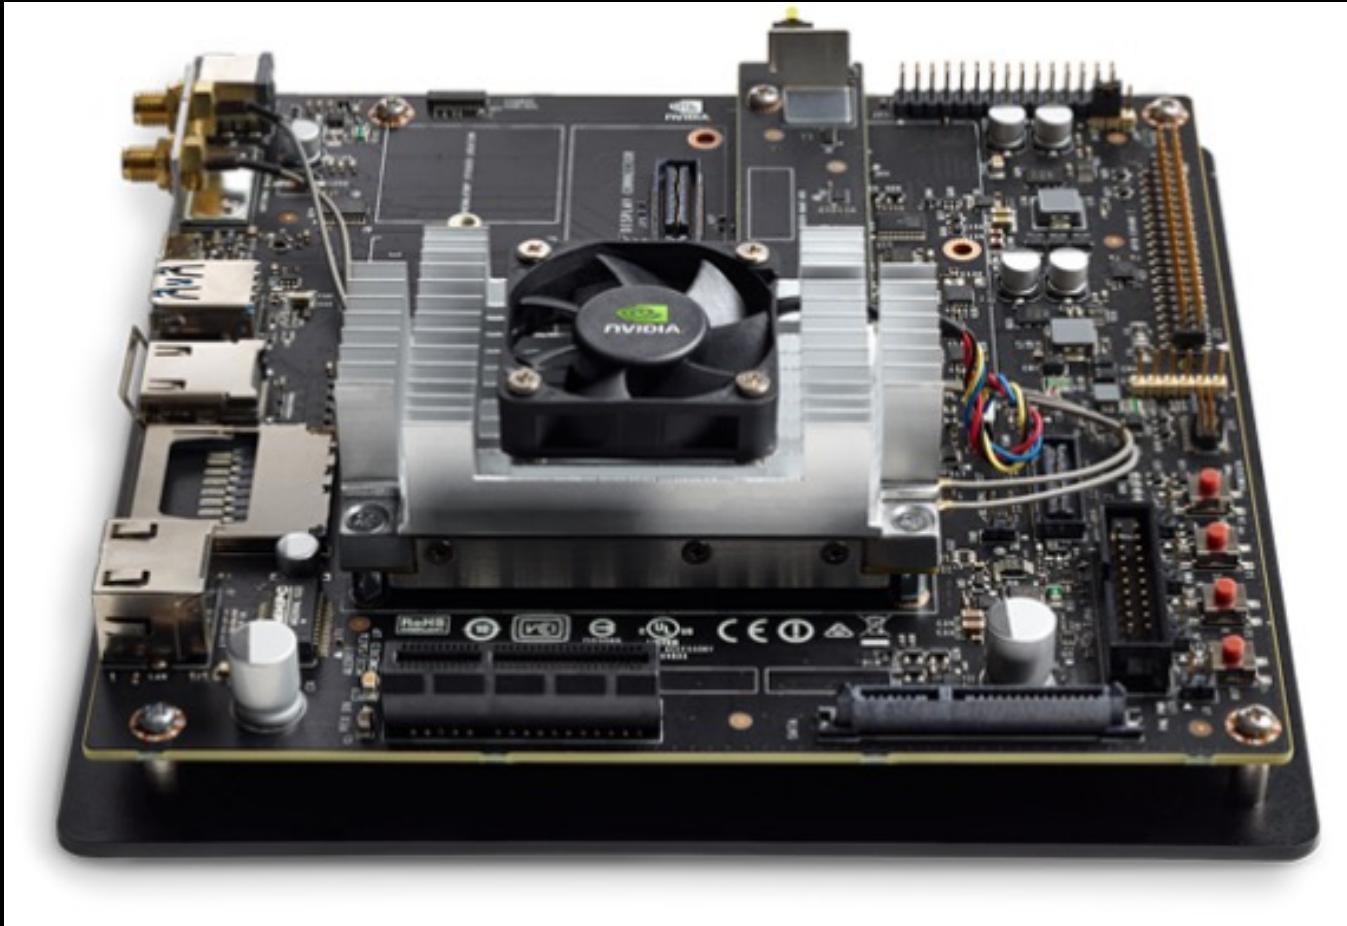
# Uncertainty maps reflect the errors



# Obligatory comparisons...



# Speeding it up for robotics



# Can use depth to improve ORB SLAM

Monocular Only

Before Loop-closure



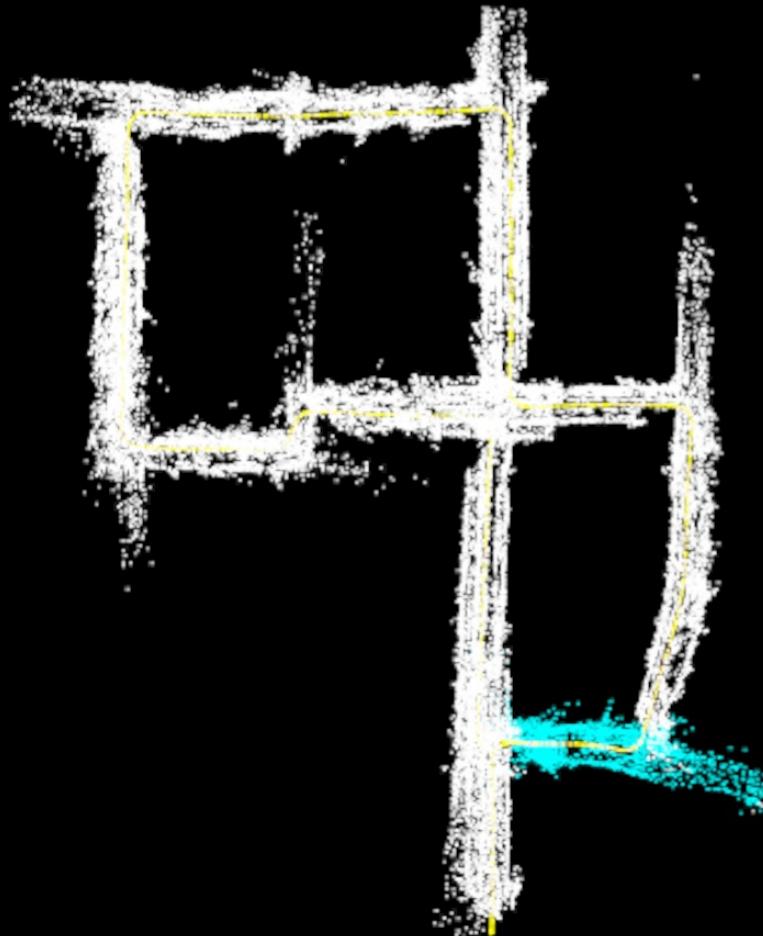
After Loop-closure



# Can use depth to improve ORB SLAM

Using Depth Predictions

Before Loop-closure

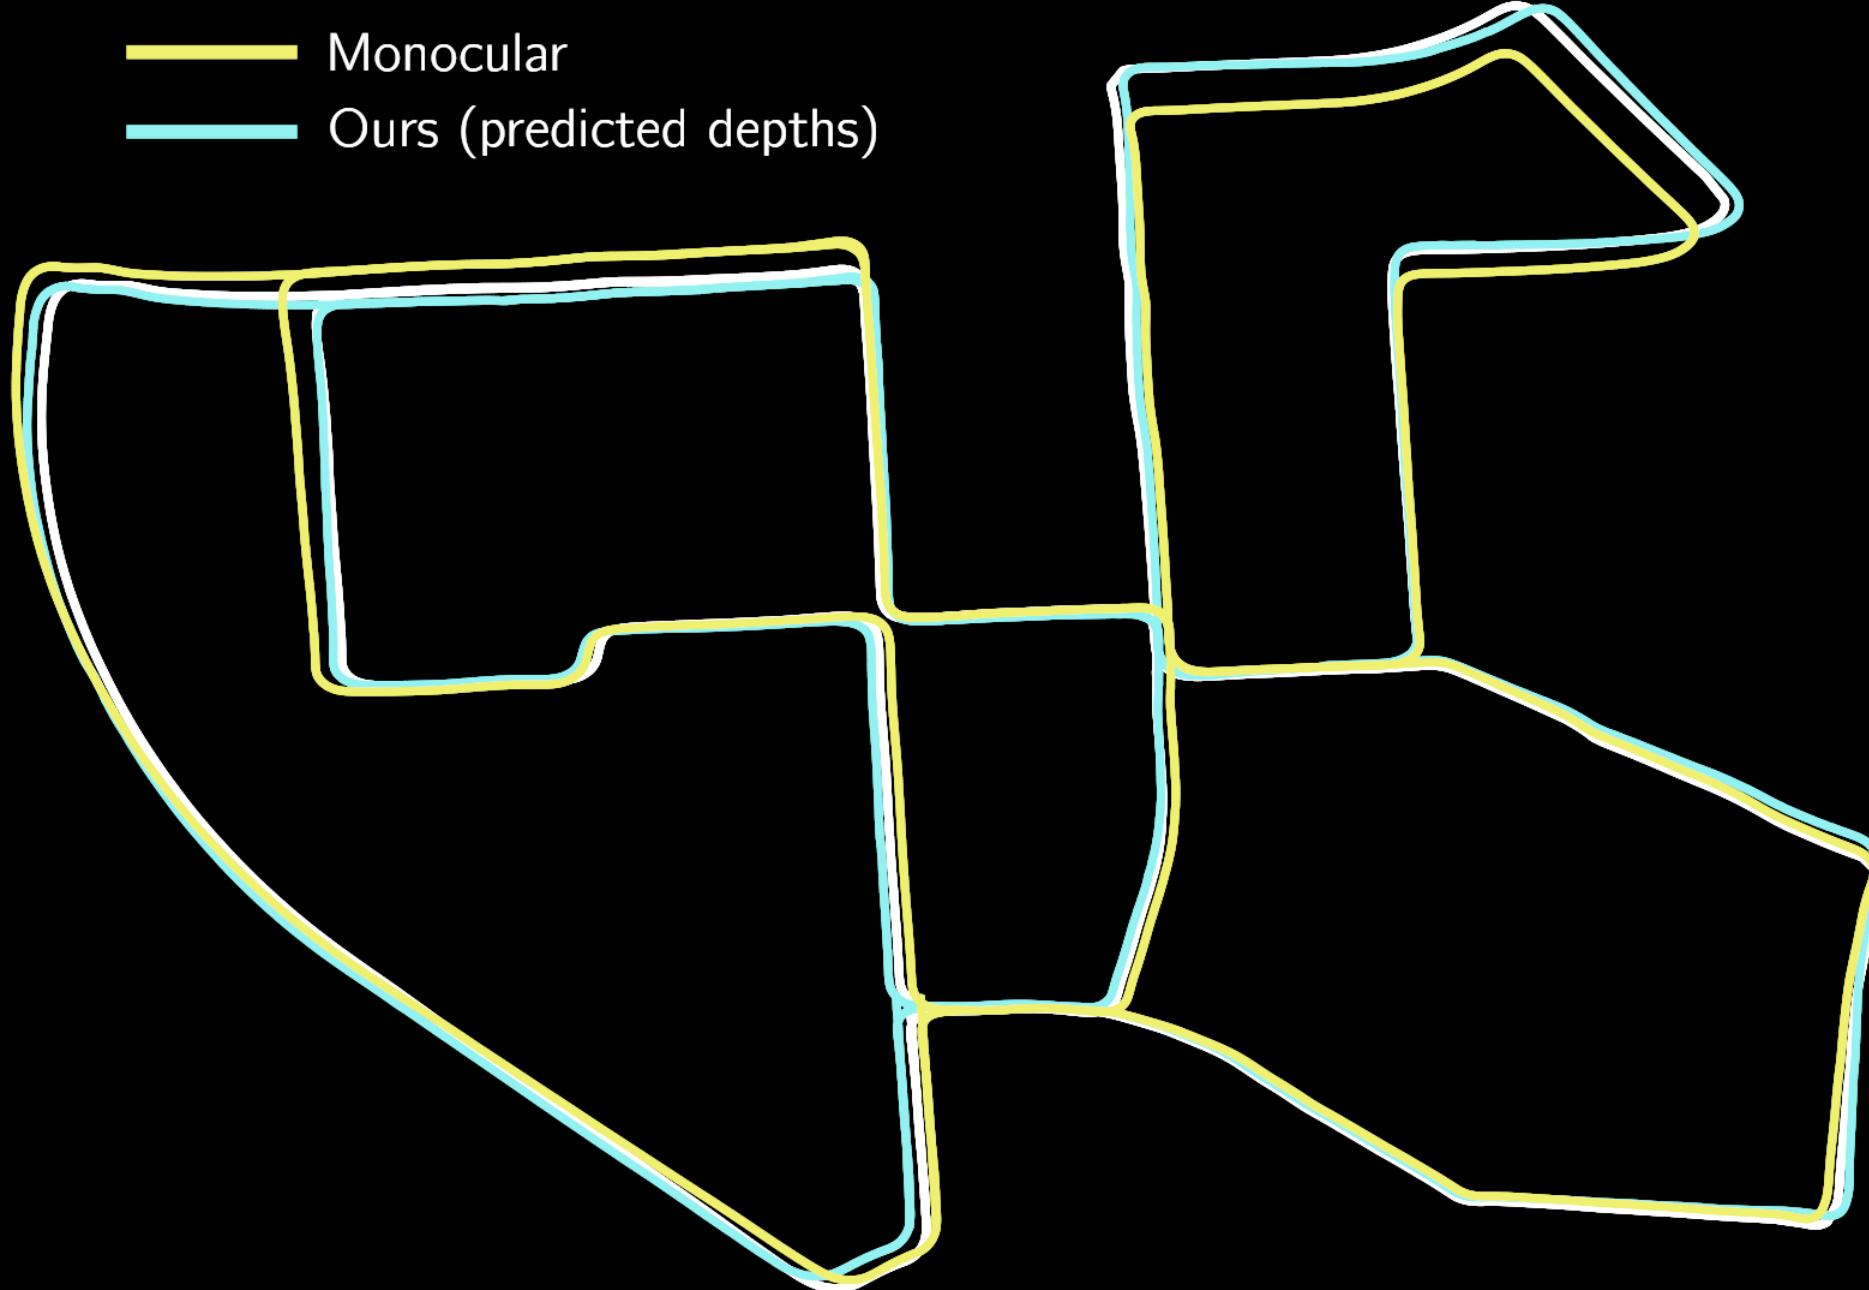


After Loop-closure



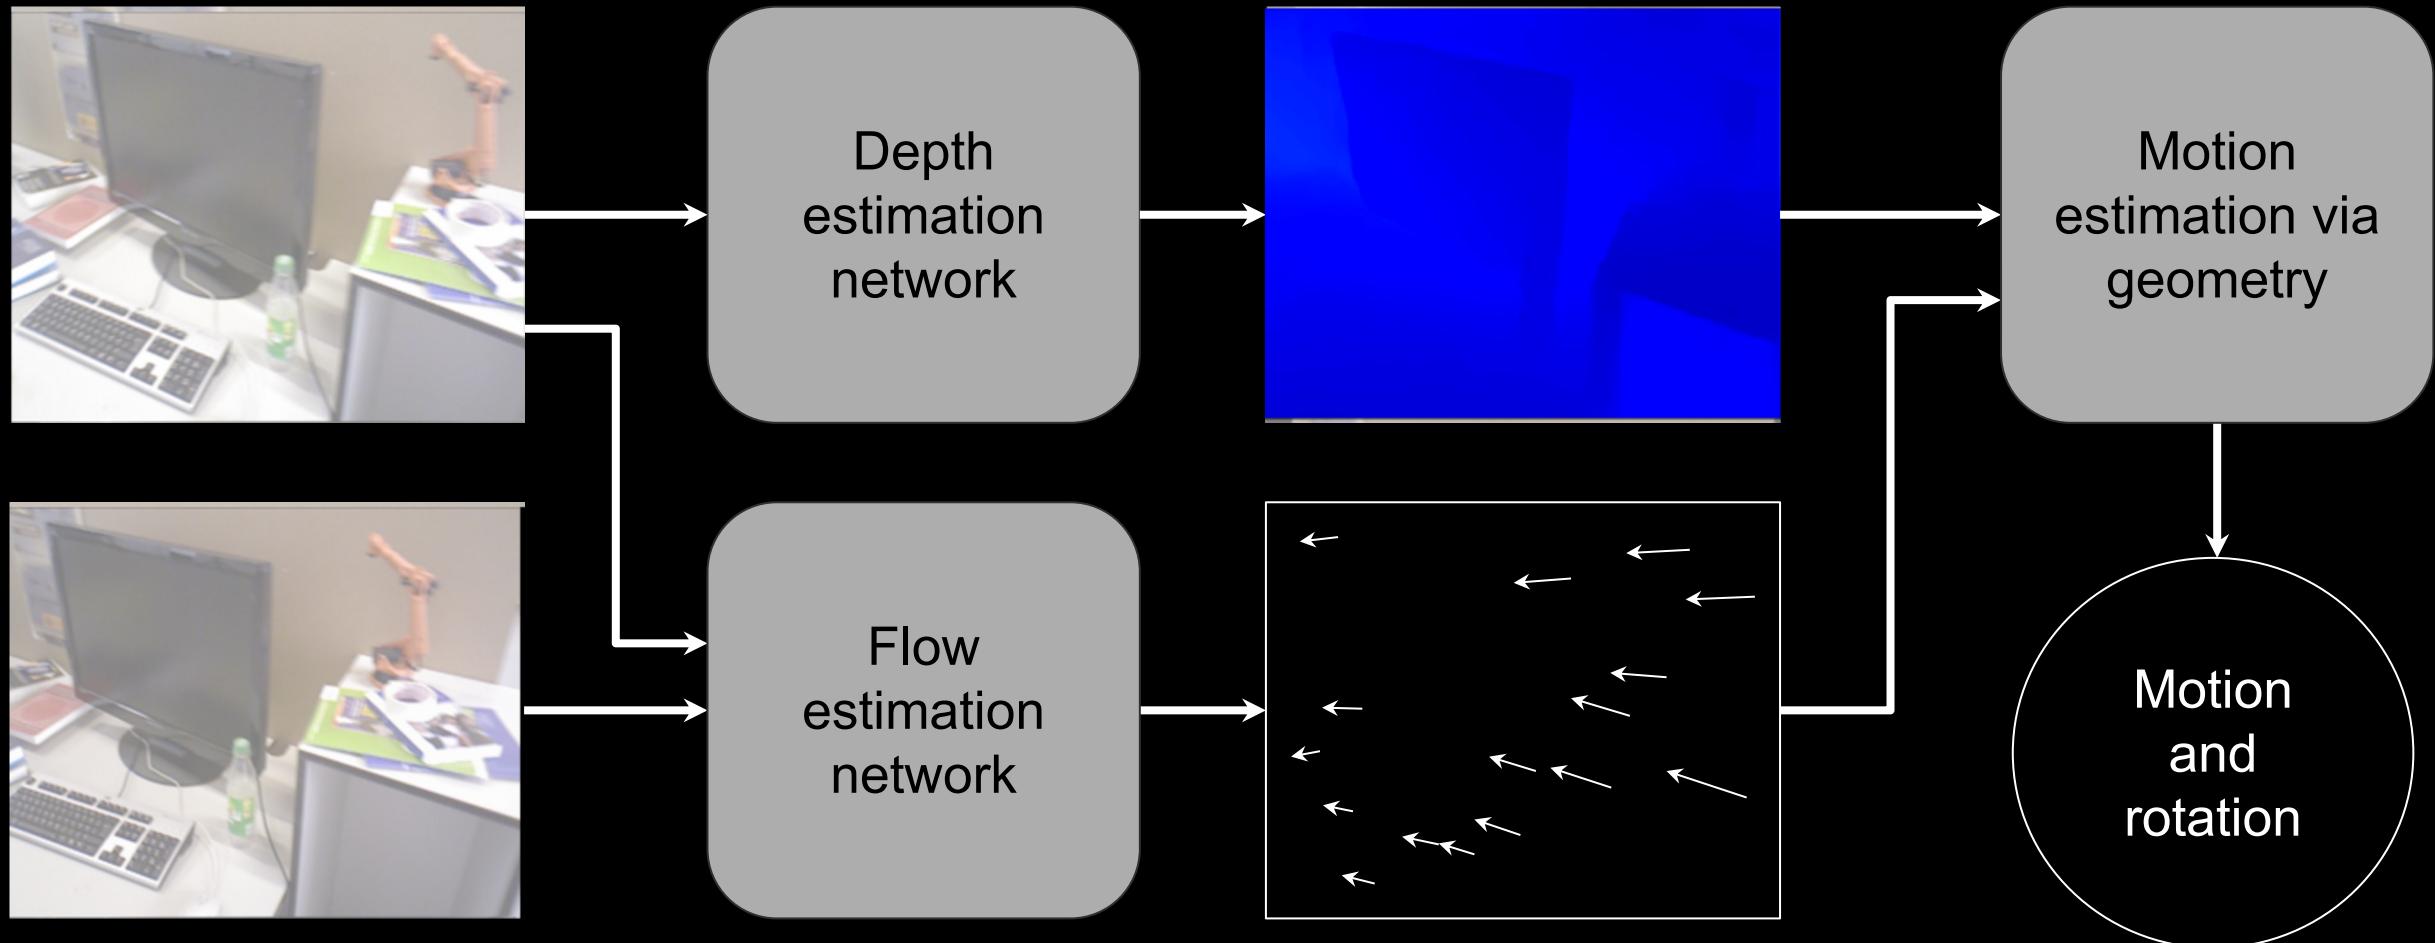
Ca

- Ground Truth
- Monocular
- Ours (predicted depths)

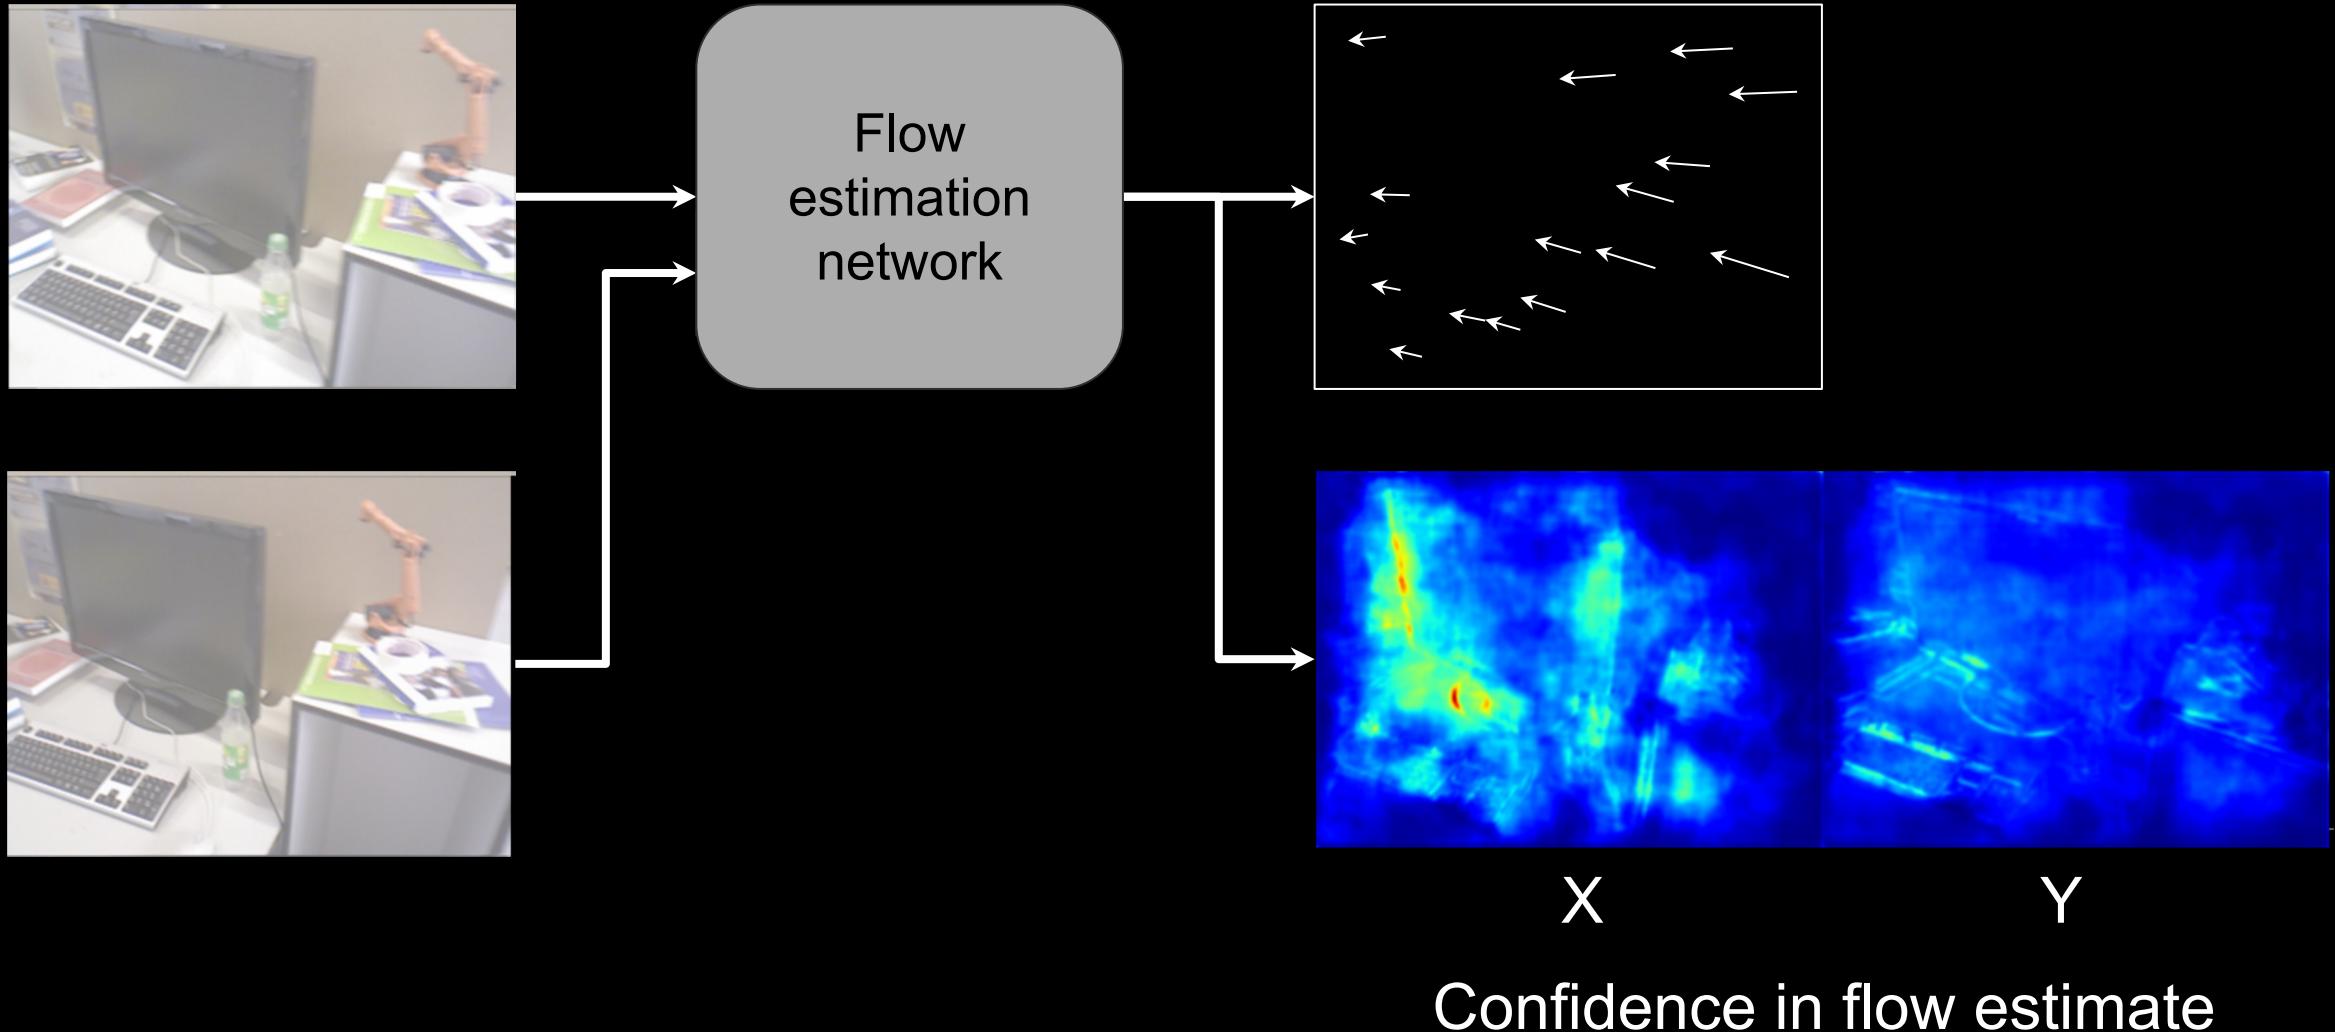


# But we want more deep learning!

Estimate flow from two images; use this to get camera motion:

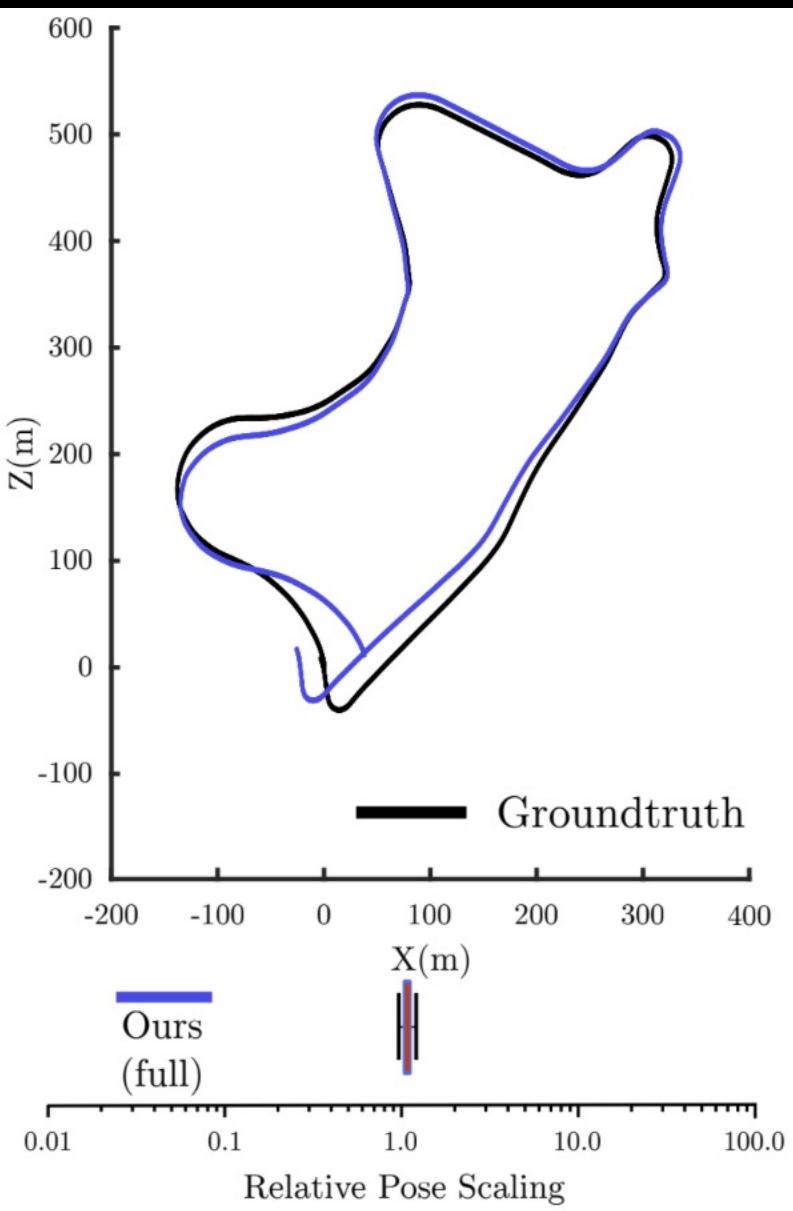
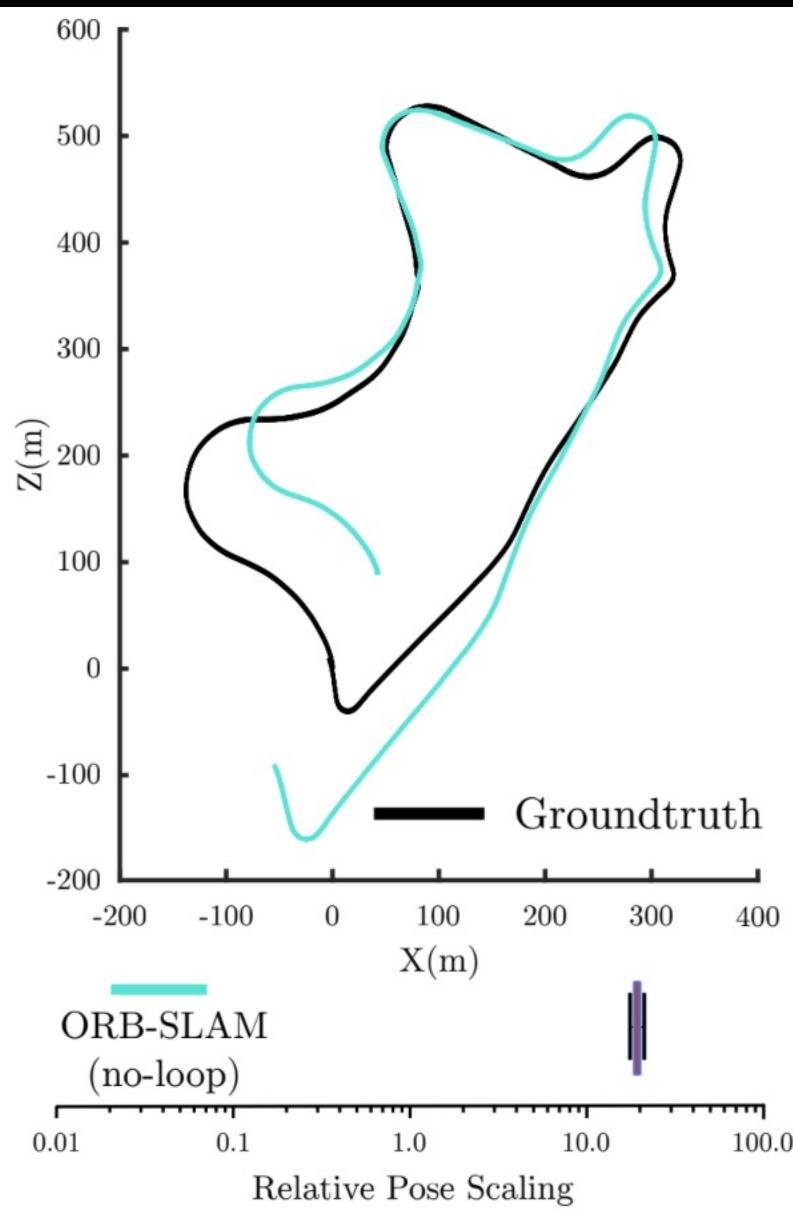
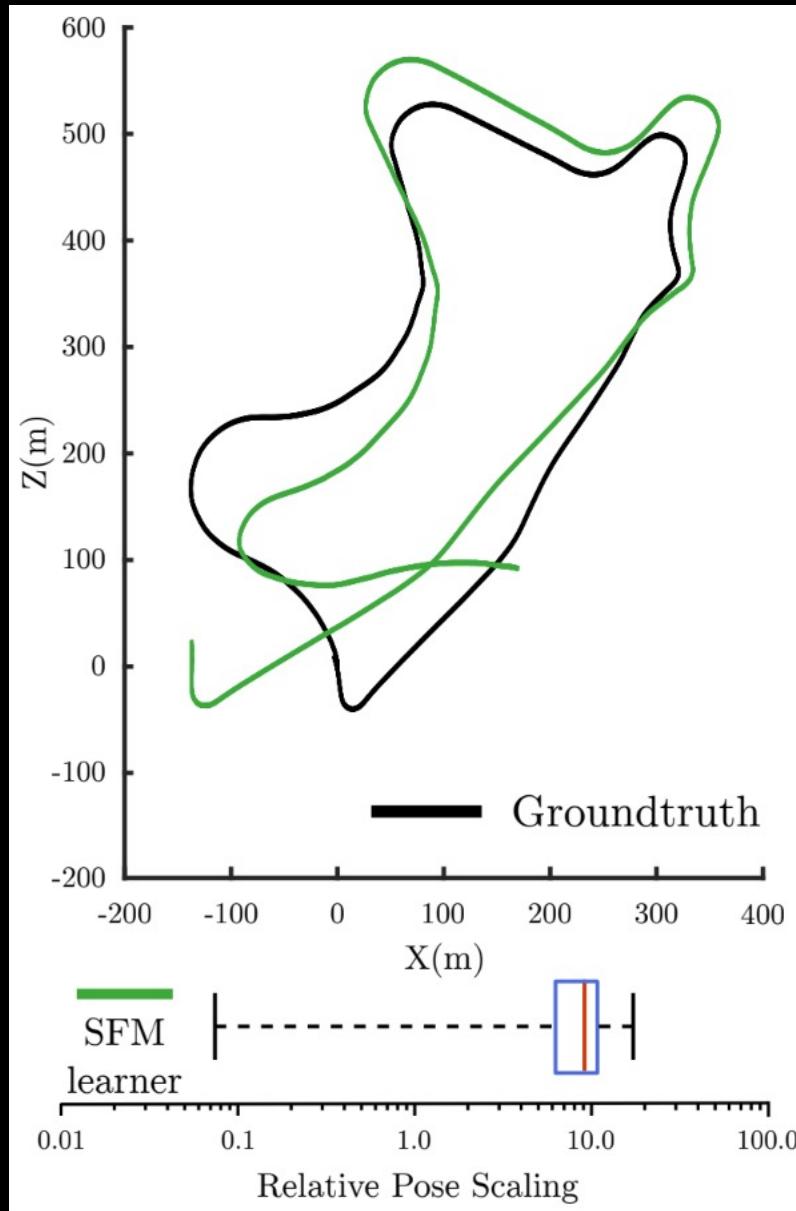


# Better if uncertainty in flow is incorporated too



# Parameterising flow uncertainty

# Visual Odometry (just prior to loop closure)



Interestingly, system learns to recognise independently moving objects and give them less confidence

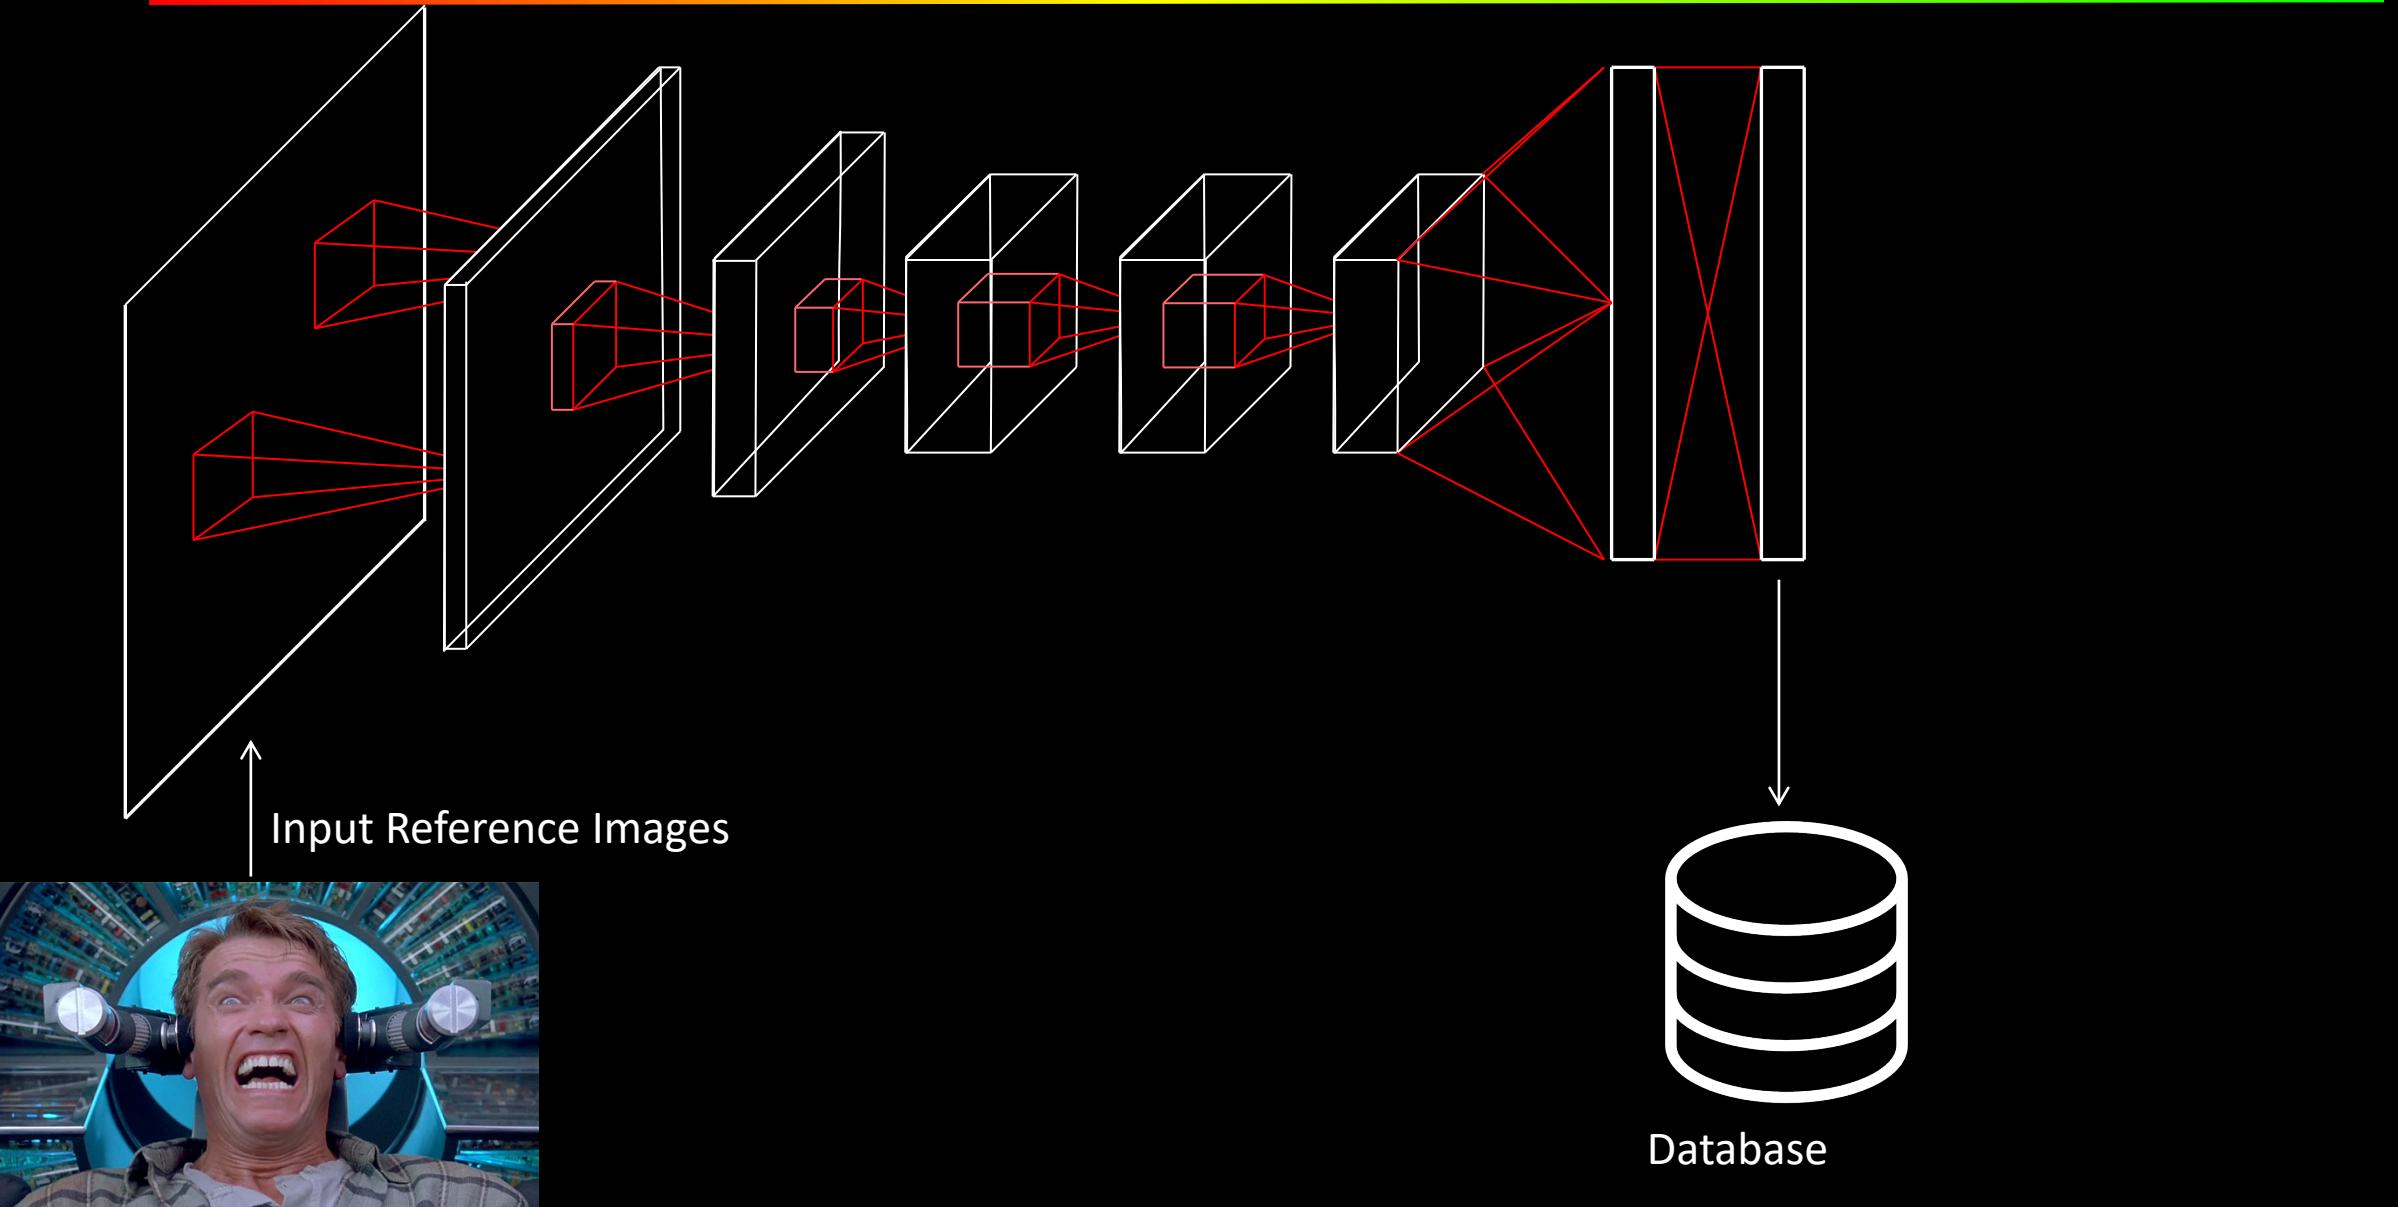


Claim 3:

You should know the difference between  
Aleatoric and Epistemic uncertainty

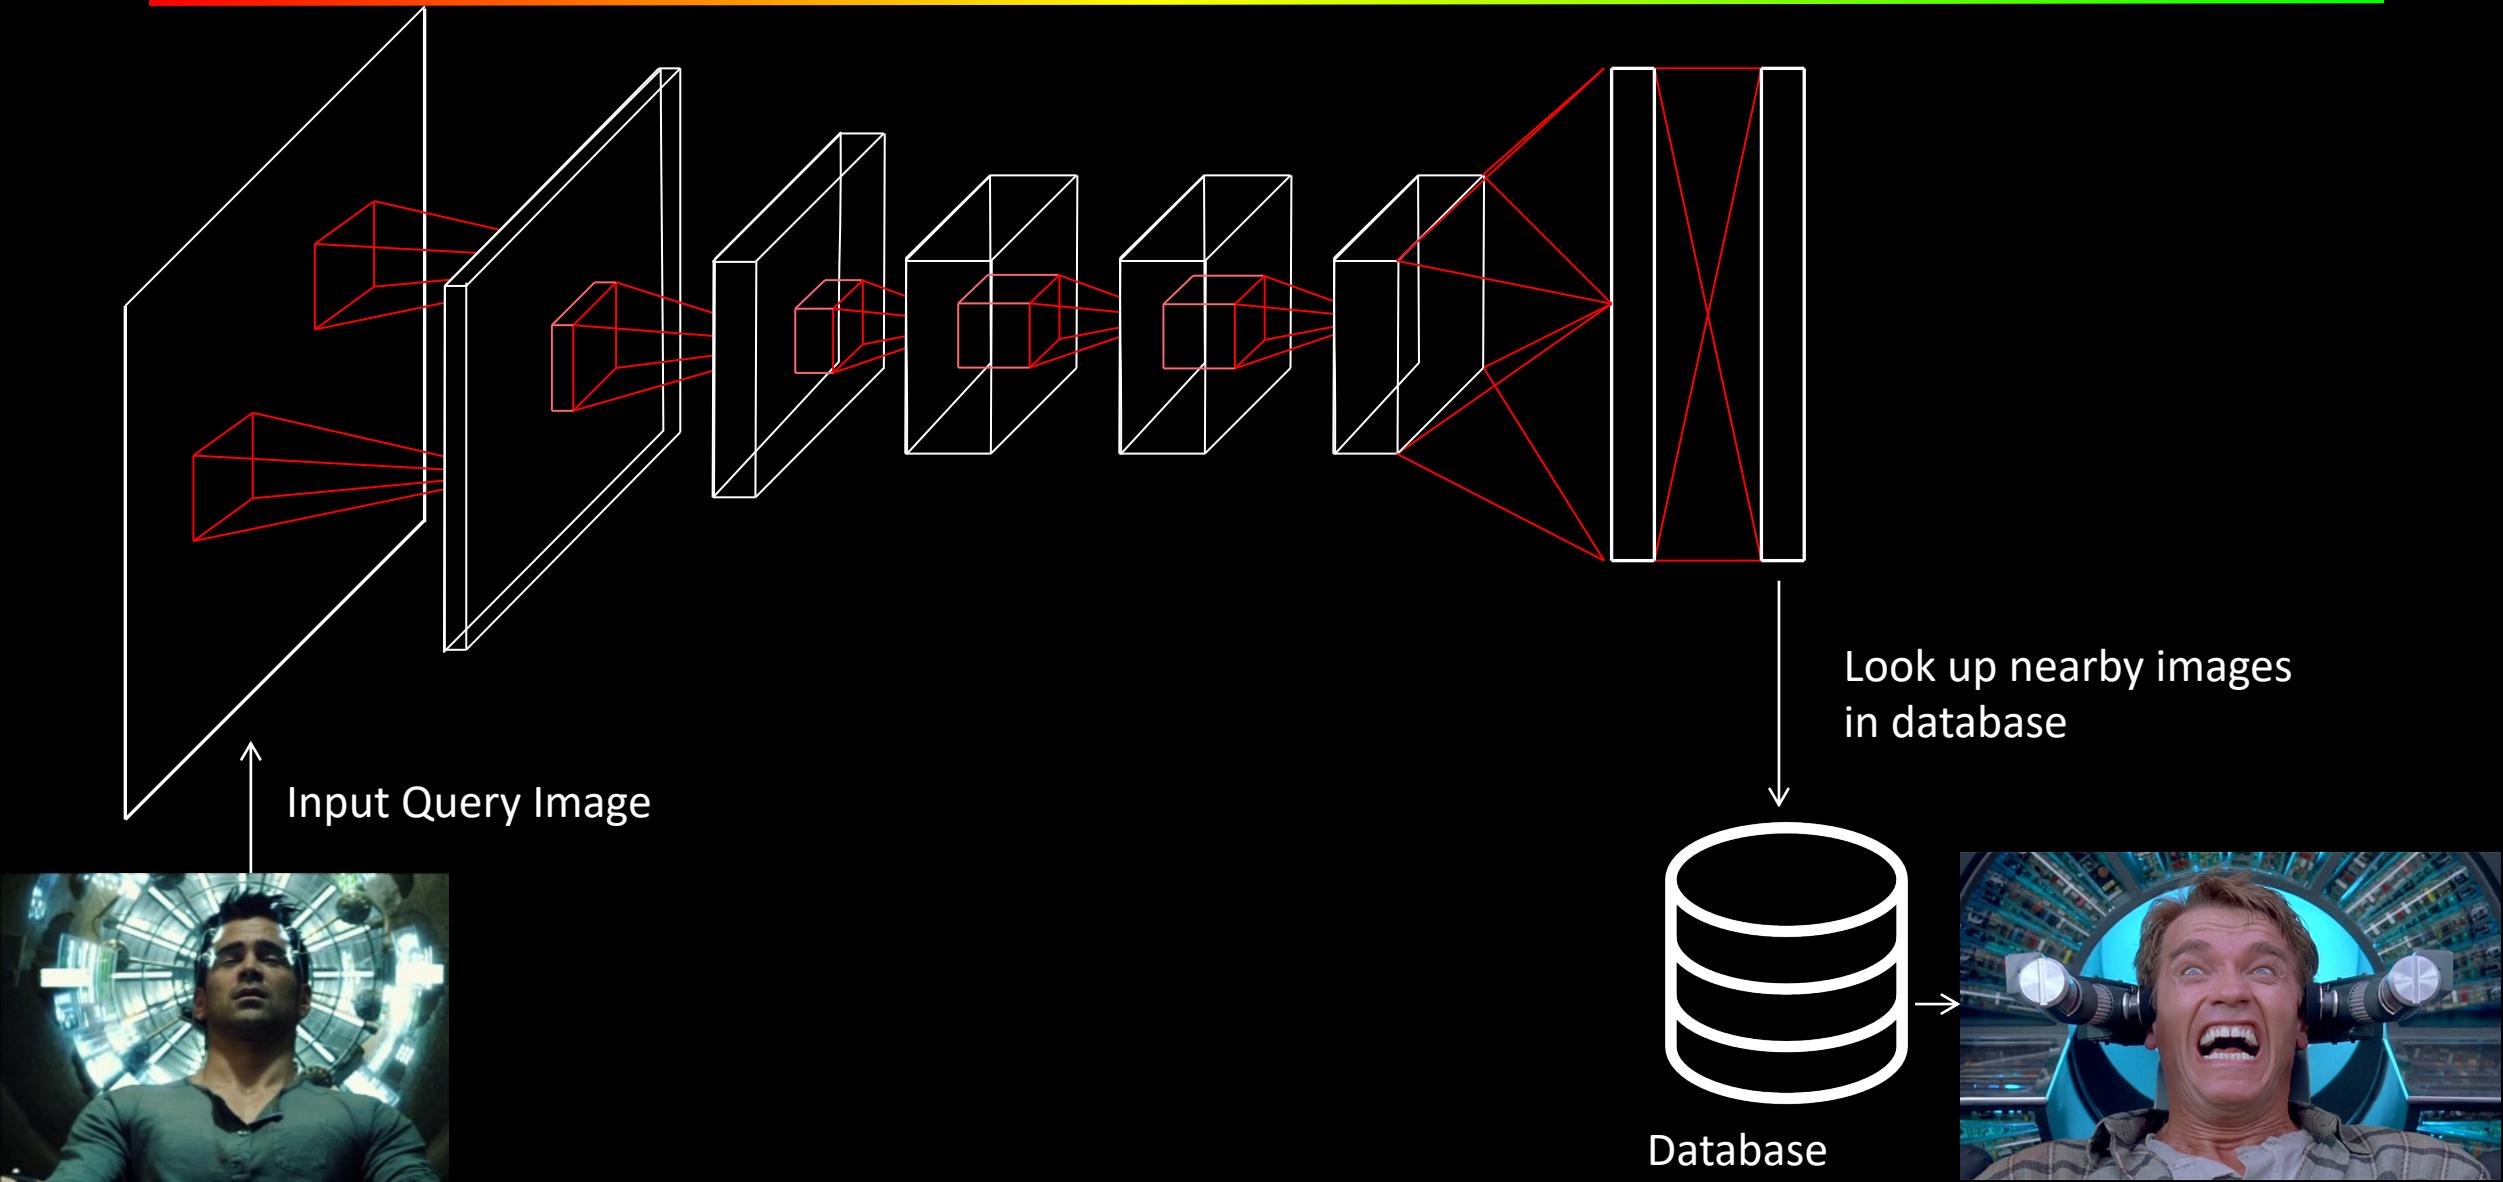
# Representation learning

---



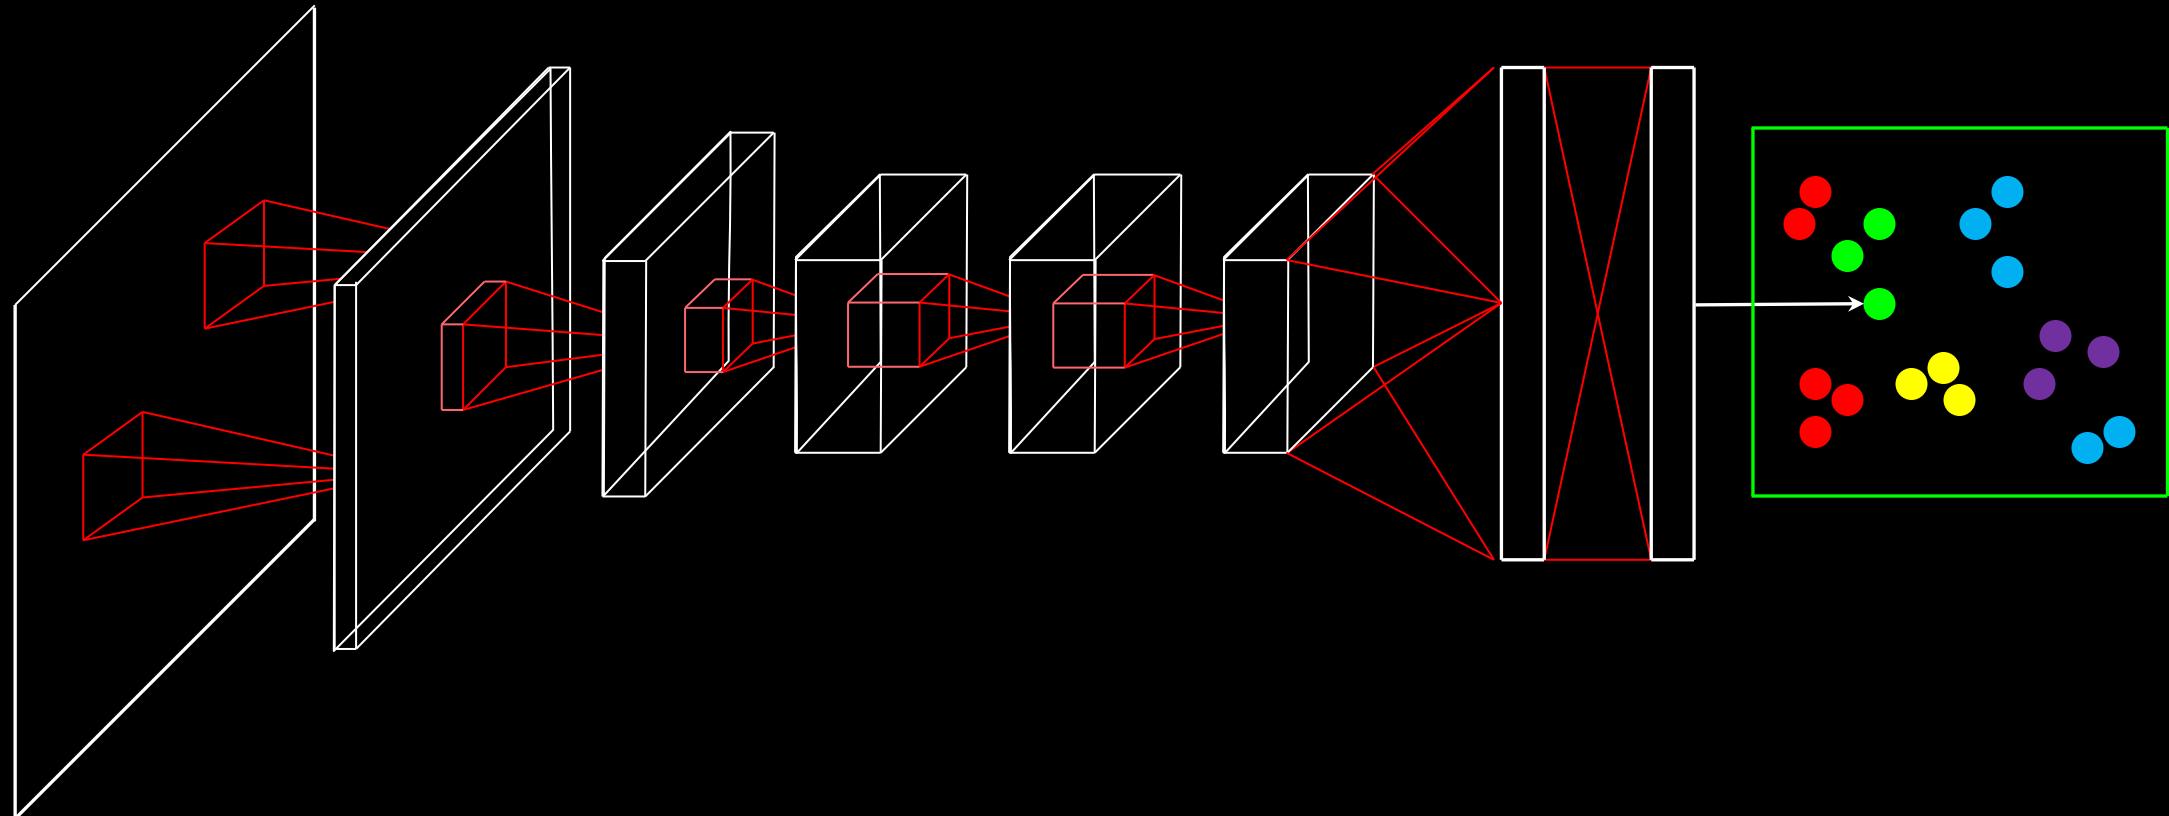
# Representation learning

---



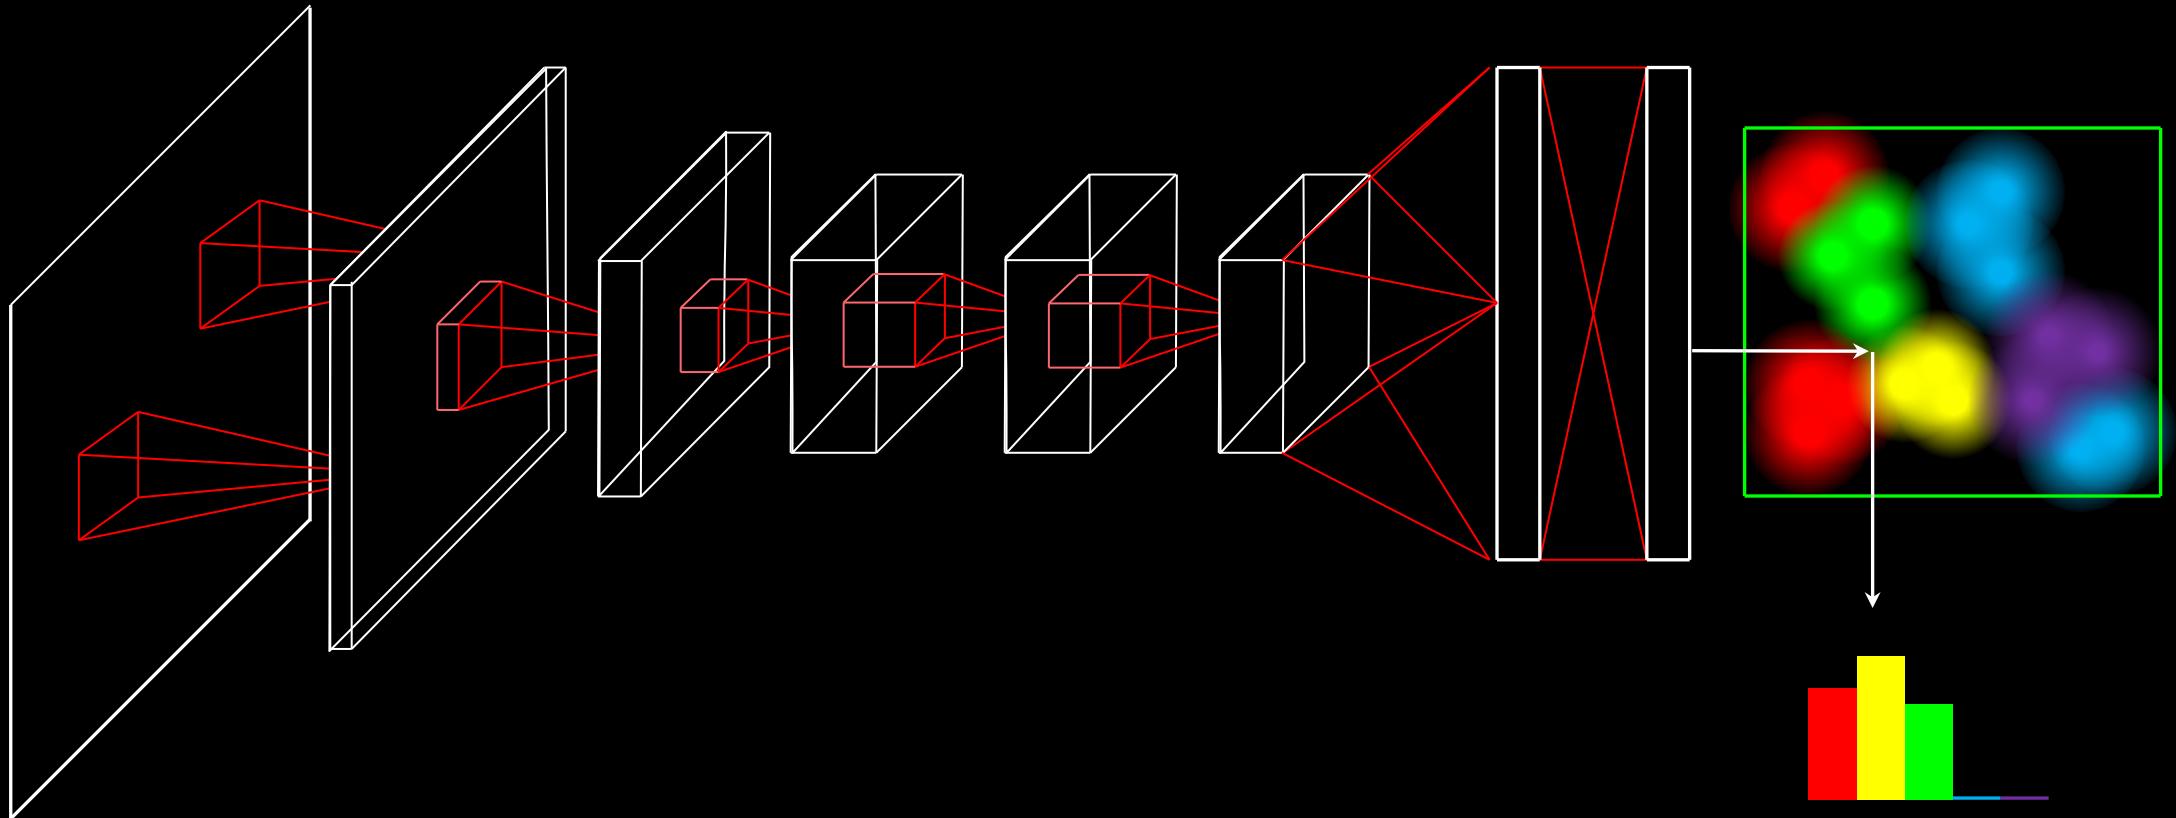
# Replace softmax with episodic memory using FANNG

---



# Use Marginal density of Gaussian kernels

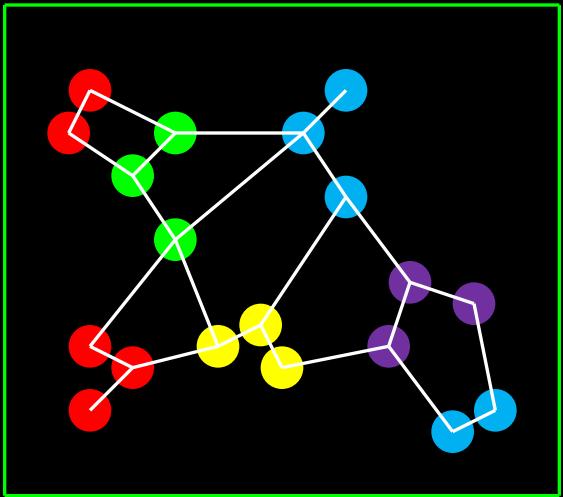
---



Cross entropy loss pushes  
examples uphill

# FANNG: Fast Approximate Nearest Neighbour Graphs

---

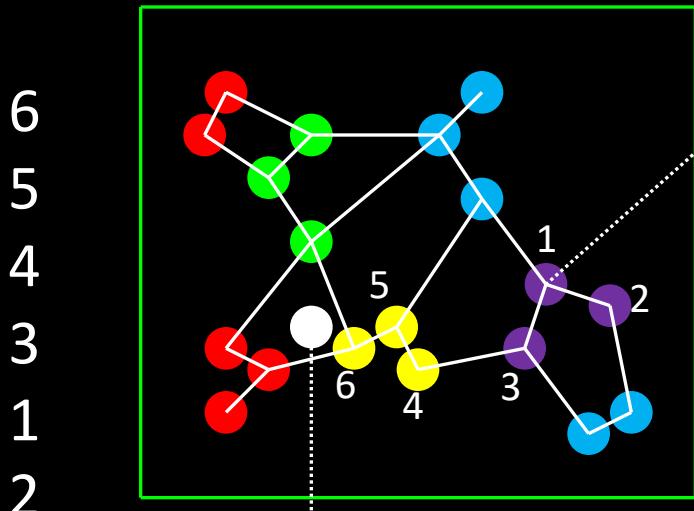


Based on two insights:

1. If a reference point is near a query point, it's worth checking its neighbors to see if they're closer
  - build a graph linking points to their neighbors
2. If there are many neighbors in the same direction it's not necessary to have edges to all of them
  - if the nearest of them isn't closer to the query then the others probably aren't either.

# FANNG Algorithm

---



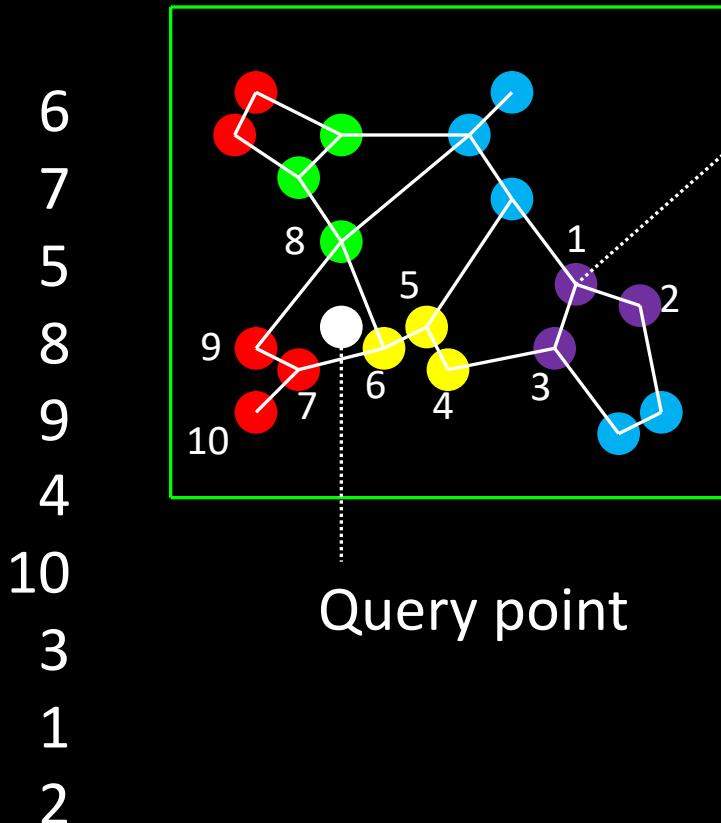
Query point

```
Start somewhere in the graph  
While (not bored) {  
    Measure distance to query  
    Insert this point into a priority queue* of vertices visited  
    Find the nearest unvisited neighbor of the highest  
    priority vertex with unvisited neighbors  
}
```

\*in order of distance from query

# FANNG Algorithm

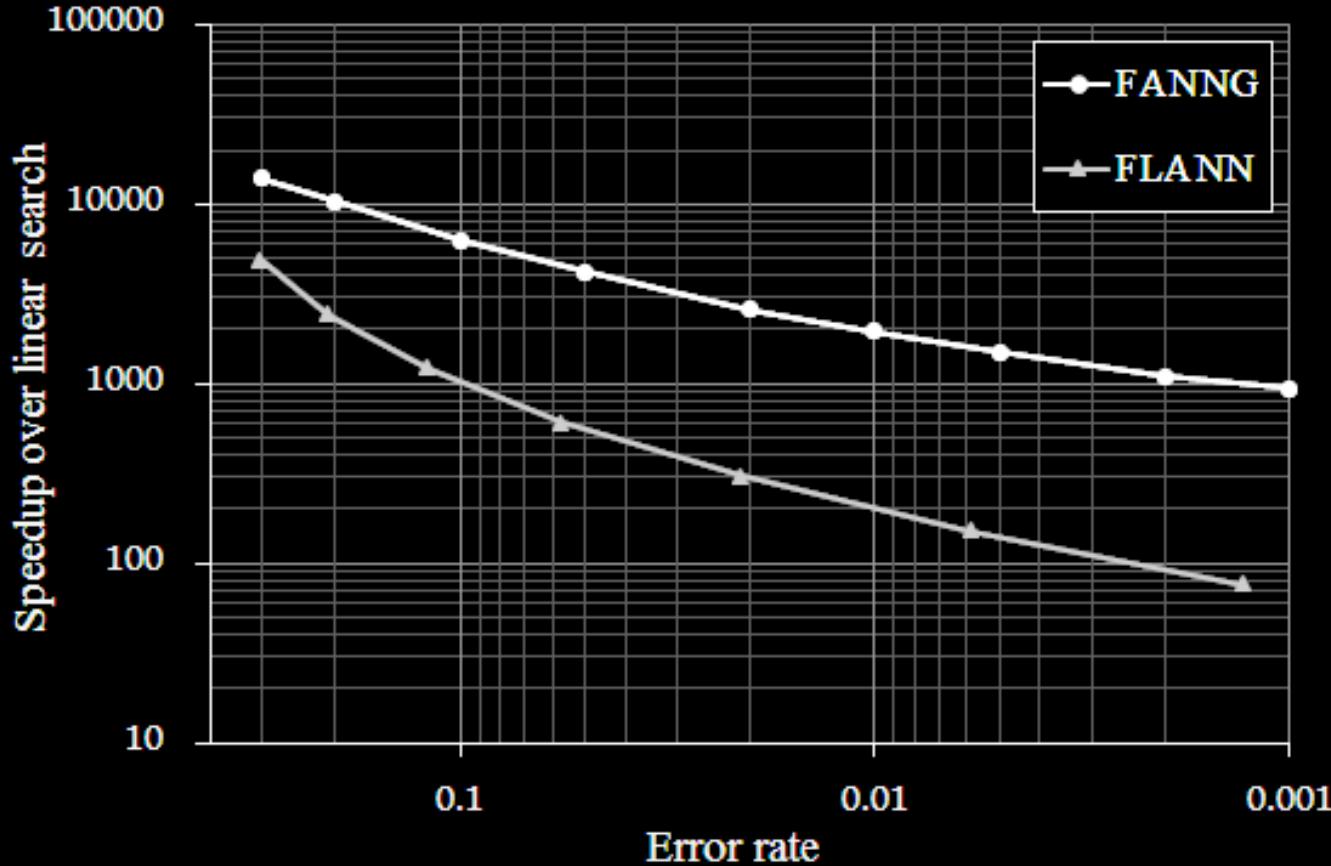
---



\*in order of distance from target

# FANNG Performance

---

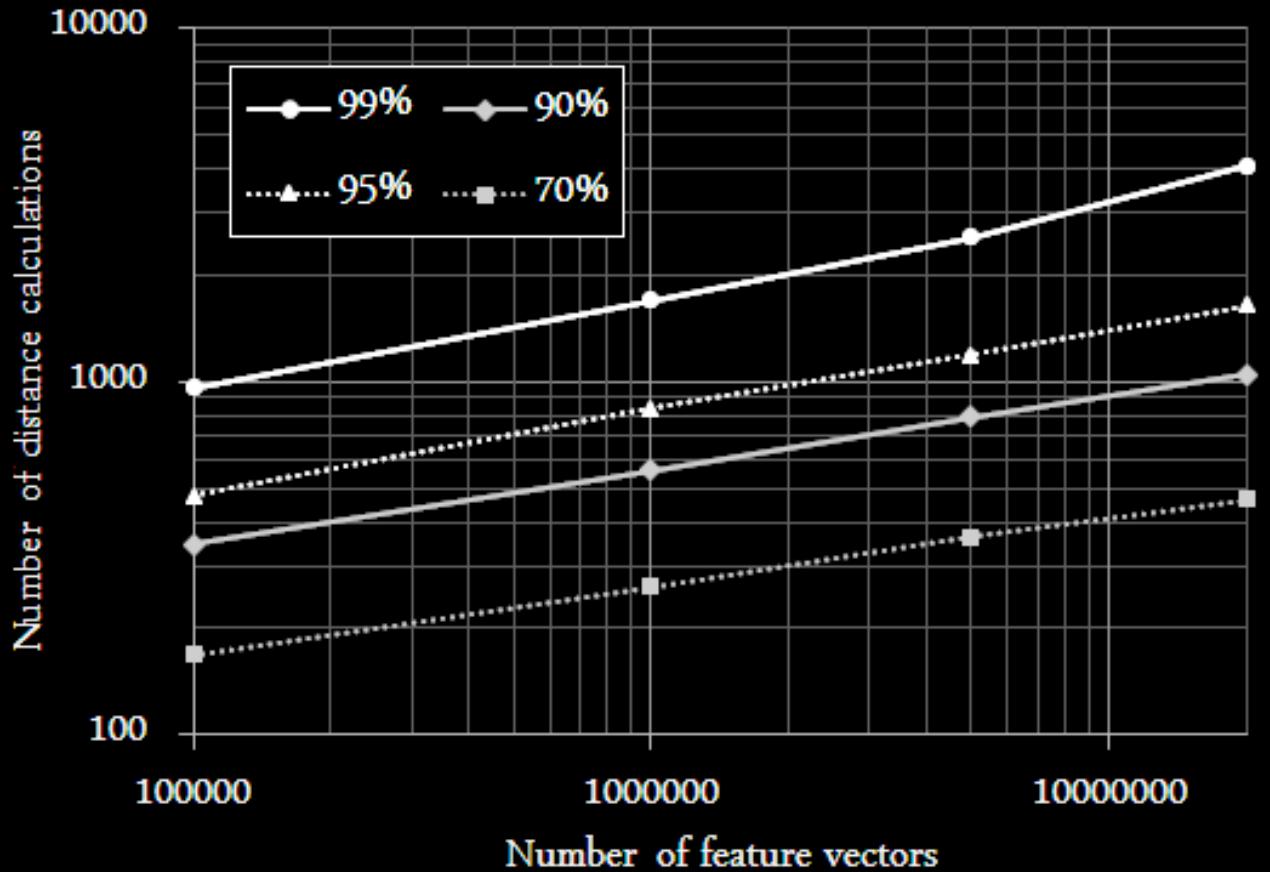


FANNG can find the nearest neighbor with 90% accuracy in 1M examples of SIFT descriptors at rate of 10M queries/s – or at 99.7% accuracy at 1M queries/s

It can also return K nearest neighbors (the highest K entries in the Priority Queue when searching is finished)

# FANNG Performance

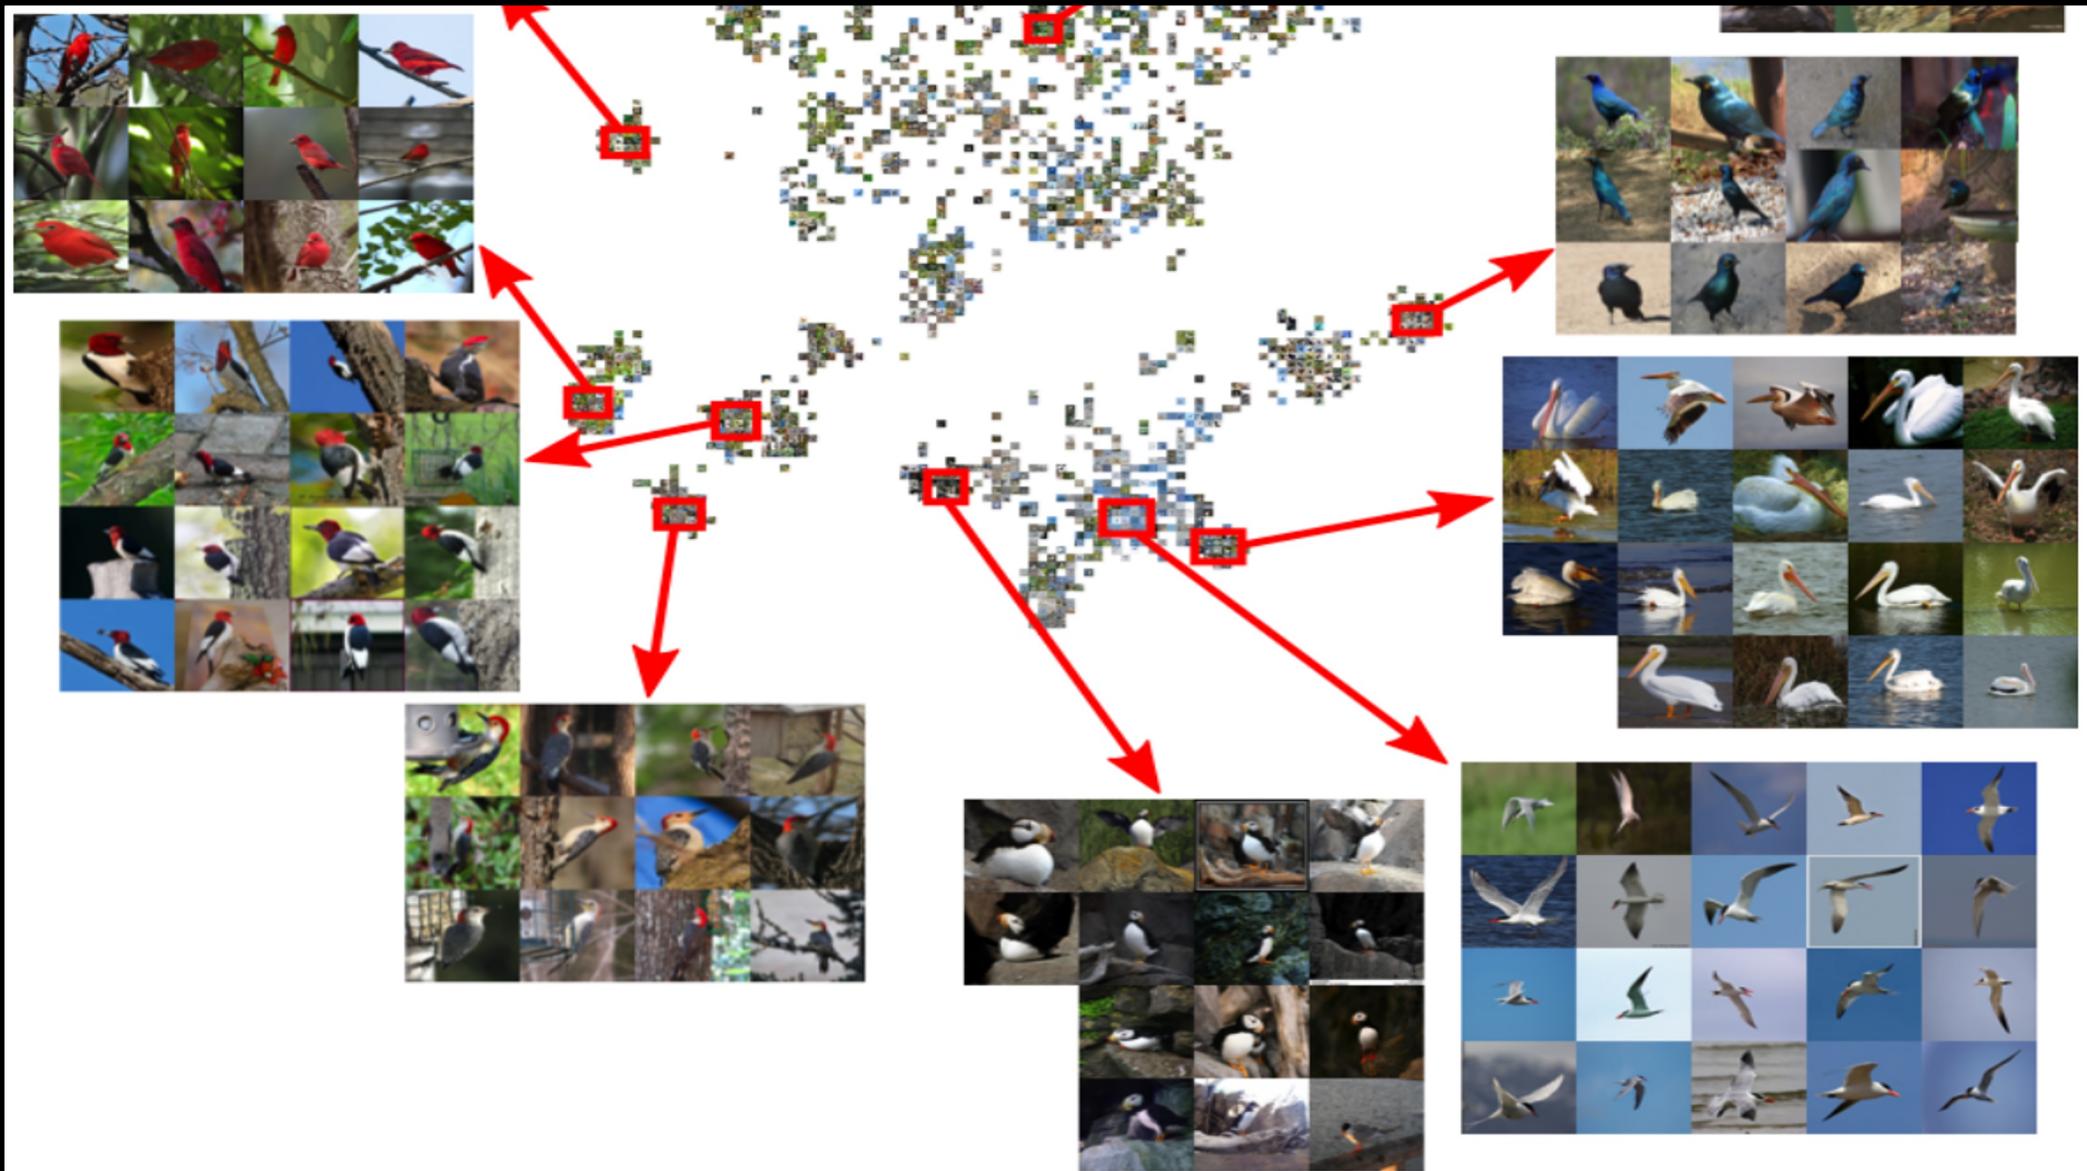
---



For 128D SIFT data  
search time scales as 0.2 power of  
data set size:

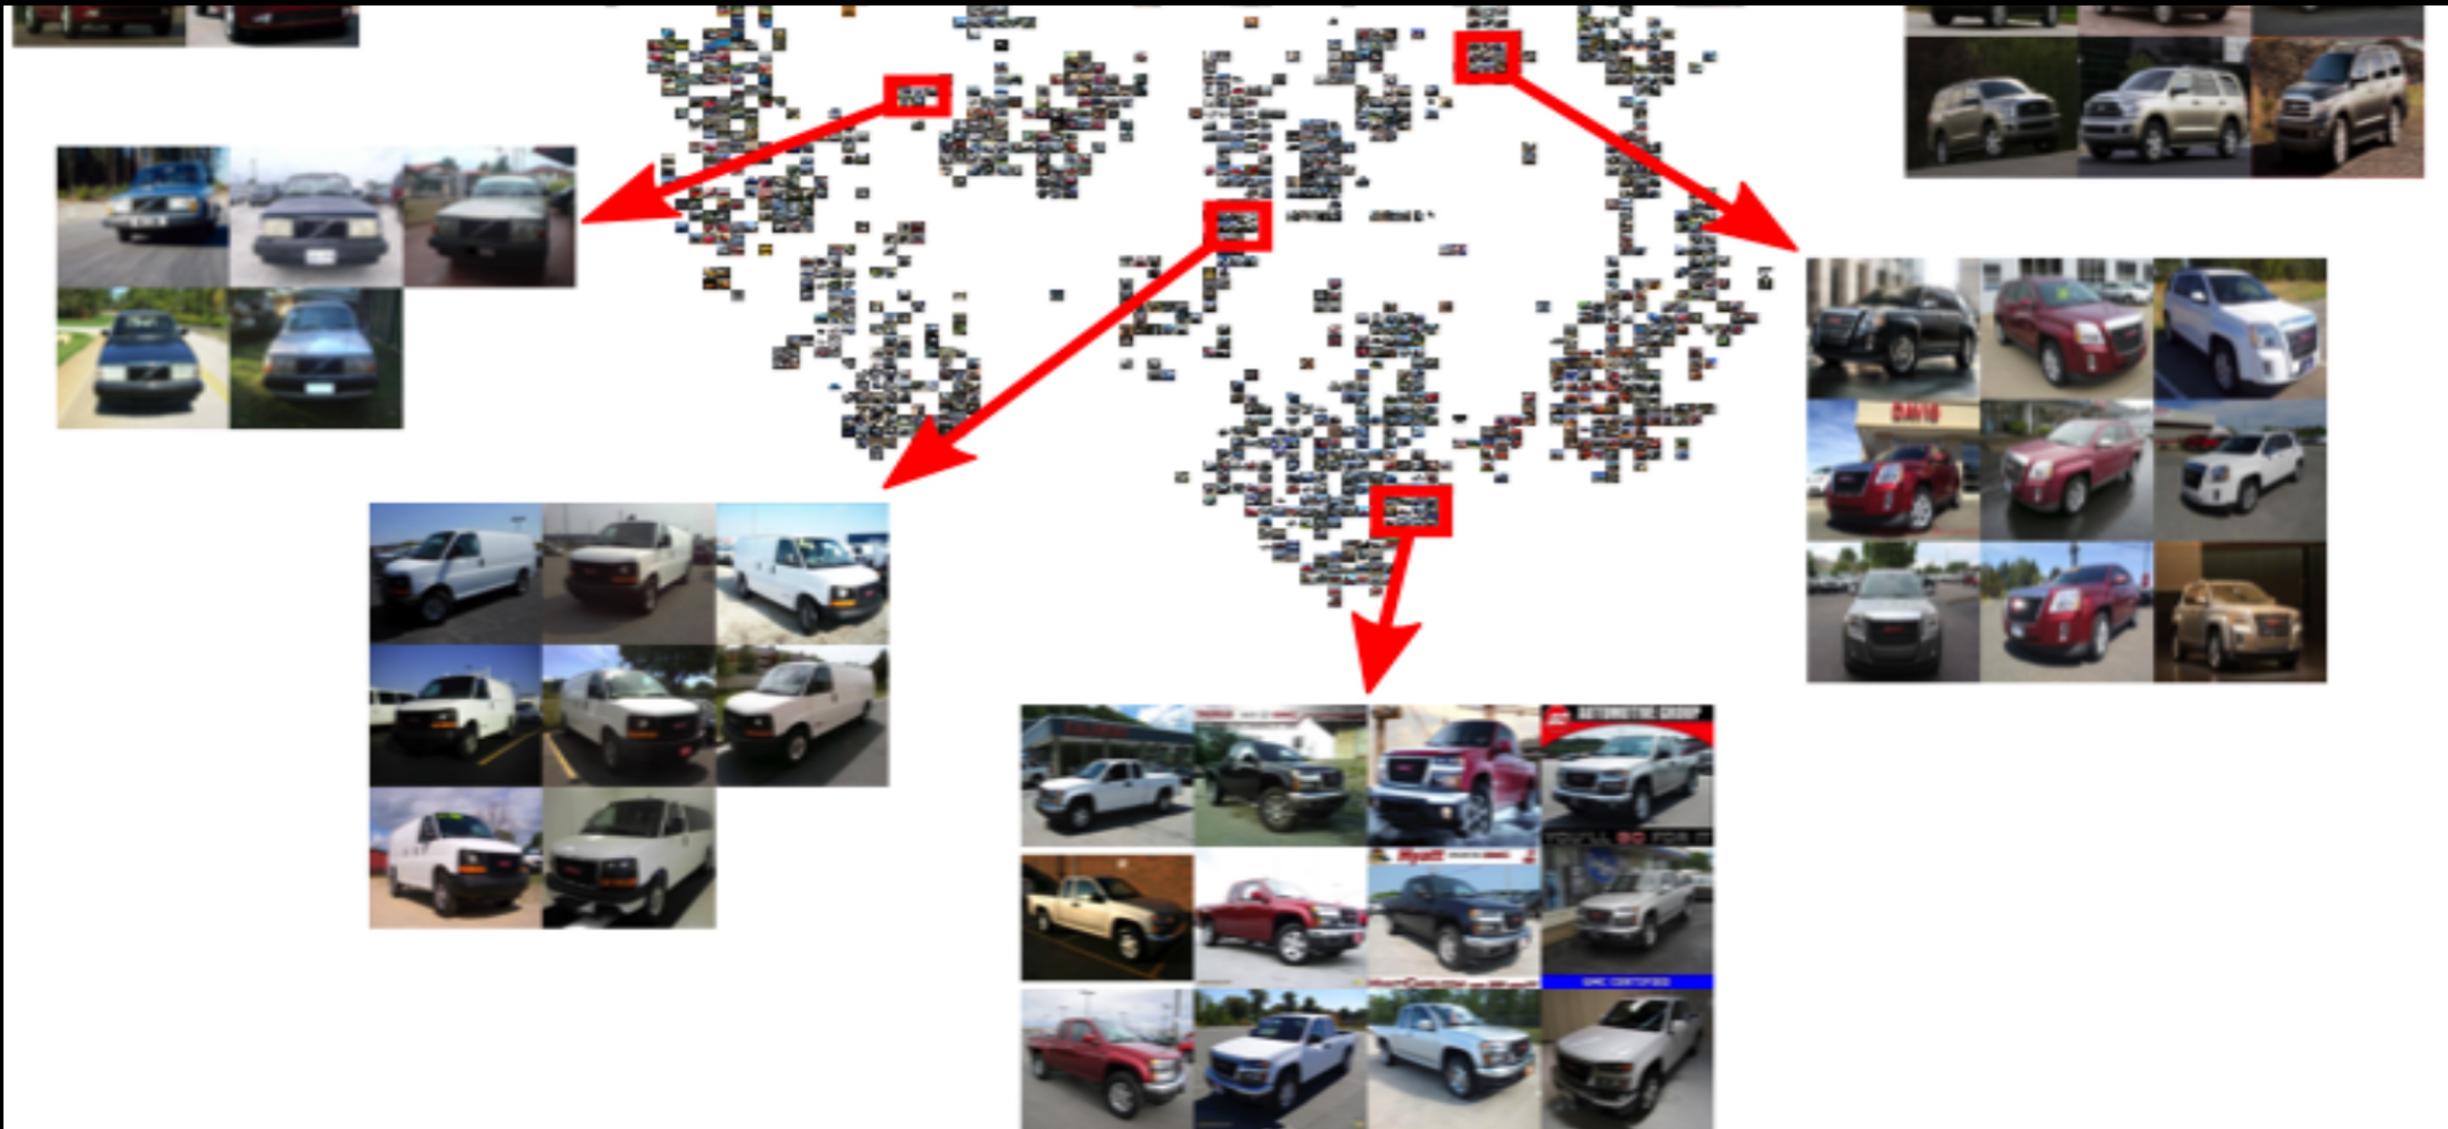
$$t \propto N^{0.2}$$

# Learned embedding transfers well



# Learned embedding transfers well

---



# Novel class detection / active learning

